


(a) a KNOWN subspace beats the transformer; LEARNING it (orange/red) does not, and even the TRUE law of supports (last bar) adds nothing



(b) more trust in the learned support is monotonely worse: the optimum is no learned support

