

The Three-Level Posterior Tower: Hierarchical Bayesian Inference and Mixed Distributive Laws

Revised version 8

Vinay K. Awasthi
Department of Applied Mathematics and Statistics
Johns Hopkins University*

October 3, 2026

Abstract

We construct a categorical framework for structured Bayesian inference that keeps three levels of posterior information as state: a posterior ρ over outcomes, a second-order posterior Π over outcome posteriors, a third-order posterior Ω , and a shadow Σ over structured hypotheses. Combining this state with the finite distribution monad D forces a design choice. In the *absolute* tower $C_3(X) = X \times DX \times D^2X \times D^3X \times \text{Sh}(X)$ the levels are re-indexed by the object, so $C_3(DX)$ already contains D^4X : the hierarchy grows by one order every time the monad is applied. We prove that $X \times E(X)$ is a comonad for every endofunctor E , that the context and distribution-level lifts are coherent by functoriality, and that the two natural candidate laws $DC_3 \Rightarrow C_3D$ each fail at least one Beck axiom; the unit axiom alone forbids point-mass hierarchy at level one. We then construct a *base-relative* tower, in which the three levels live over a fixed outcome base B and do not grow. Its *coherent* contexts, those whose barycentric projections and shadow predictive agree, form a sub-algebra of the monad, and the canonical hierarchy $\rho_h, \Pi_h = \delta_{\rho_h}, \Omega_h = \delta_{\Pi_h}$ extends to an affine section of it. For this tower we obtain a strict mixed distributive law satisfying all four Beck axioms, natural in the base. Pooling commutes with hierarchy formation, conditioning commutes with pooling after evidence reweighting, and distinct posteriors are retained at level two. We then pose a coherent lifting problem: given admissible first- and second-order states generated by a finite hypothesis class, when does a claimed posterior state extend to a full coherent three-level context? Admissibility at each level is a convex condition, failure is monotone up the tower, each level can be the first to fail, and every failure carries an explicit certificate: a separating test function at level one (with the obstruction size given by a minimax identity) and a witness state at levels two and three. Pooling can hide level-one obstructions but never level-two or level-three ones. When the admissible second-order states are restricted, a claim can pass every local test and still admit no coherent lift; we characterize such claims by effective admissible sets with explicit certificates, including second-order test functions on posteriors, and define a secondary obstruction on the fibre of level-two lifts whose value depends on the chosen lift and ranges over an interval that contains zero exactly when a lift exists. This gives a stratified, precise candidate definition of a hallucinated claim as one whose required coherent lift is obstructed. We also characterize when a mixed context can be attributed to its sources: weights are unique exactly for affinely independent sources, the supports-only case gives a sharp belief-plausibility interval, and under contamination by unmodeled sources the data certify a lower bound on the unmodeled weight, while every exclusion depends on an assumed budget and the number of unmodeled sources is not identifiable. For generated

*Code and verification scripts: https://github.com/vkawasth/Trainingless-Transformer-NoGradientDescent/tree/main/Heirarchical_posteriors

text we give a likelihood projection onto two declared contexts and a reference, with a profile band for the share of each context, a certified lower bound on the share explained by neither, and a level-two test that separates blending the contexts from switching between them; folding the reference into the context models makes the unexplained share unidentifiable, and on real news text with known truth the projection recovers context shares within its bands and reports text from outside the context as unknown. Analytic sieving and its parity barrier serve as a classical illustration of iterated conditioning and observational equivalence. We identify the coalgebras of the base-relative comonad with sets over the context algebra and its bialgebras with convex spaces over it, and we make precise that the ambient 2-category is **Cat**, in which these categories are locally discrete, so two-dimensional monad theory adds no content here. We then test the natural relaxation of the absolute tower by the convex order: the unit and multiplication defects of the shift law are comparable but in opposite orientations relative to the Beck diagrams, and the comultiplication defect is incomparable, for both candidate laws. The naive convex-order relaxation therefore fails, which motivates an enrichment that compares coherent lifts rather than barycentric spreads. On the algebraic side we give a normalized partial composition on structured posteriors, corrected conditioning laws, an exact characterization of when observational equivalence is a congruence, and a proof that the likelihood-factorization “iff” holds under nondegeneracy alone. All identities and counterexamples were checked in exact rational arithmetic; the scripts are in the accompanying repository.

Contents

1	Introduction	4
2	Preliminaries	5
2.1	The finite distribution monad	5
2.2	Mixed distributive laws	6
3	Structured hypotheses	6
4	The structured posterior	7
4.1	Mixture	7
4.2	Composition	7
4.3	Conditioning	8
5	Observational equivalence	8
5.1	Equivalence and congruence	8
5.2	Identifiability of decompositions	9
5.3	Robust decomposition under unmodeled contamination	10
5.4	Sequential attribution of generated tokens	12
6	The three-level tower and the context comonad	15
6.1	The three posterior levels	15
6.2	A general context comonad	15
6.3	The three-level context	16
6.4	The context lift	16
7	The distribution-level lift	17

8	Option 1: mixed distributive laws for the absolute tower	17
8.1	What the unit axiom forces	17
8.2	The shift law	17
8.3	The averaging law	18
9	Option 2: the base-relative three-level hierarchy	19
9.1	Coherent contexts	19
9.2	The canonical hierarchy	20
9.3	The base-relative comonad	20
9.4	The strict mixed distributive law	20
9.5	Pooling, hierarchy and conditioning	21
9.6	Change of base	22
9.7	Comparison of the two options	22
10	Coherent lifting and obstruction certificates	22
10.1	Admissibility data and the admissible tower	22
10.2	Claims, lifts and failure levels	23
10.3	Obstruction certificates	24
10.4	Effective admissibility and secondary obstructions	25
10.5	Interpretation and scope	27
10.6	Sieving as iterated conditioning: a classical illustration	28
11	Nondegeneracy and the likelihood-factorization criterion	28
12	Coalgebras, bialgebras and the 2-categorical setting	29
12.1	The ambient 2-category	29
12.2	Coalgebras of the base-relative comonad	29
12.3	The lifted monad	30
12.4	Bialgebras	30
12.5	What two-dimensional monad theory does and does not add	31
12.6	Convex-order Beck comparisons for the absolute tower	31
13	Empirical setup	33
14	Code and data availability	33
15	Related work	33
16	Summary of results	34
17	Open problems	35
18	Conclusion	36
A	Empirical tables	37
A.1	The 40-world benchmark	37
A.2	Chain-pooling curve	37
A.3	Coverage	37
A.4	Shallow-MLP worlds	37
A.5	Spectral pre-pass	37

1 Introduction

Bayesian inference in structured problems produces a hierarchy of posteriors: a posterior over outcomes, a posterior over hypotheses (each of which itself predicts outcomes), and a posterior over the ways hypotheses combine. This paper builds a categorical framework that retains all three levels as state rather than as intermediate computations, and asks how such state interacts with the probability monad.

The framework builds on the three-layer architecture of [5]; the hierarchy layer is developed in [3] and the holonomy obstruction in [4].

The design choice. The levels of the tower are indexed by the object on which the comonad is evaluated. Evaluating at DX , as any distributive law $DC_3 \Rightarrow C_3D$ must, shifts every level up by one order:

$$C_3(DX) = DX \times D^2X \times D^3X \times D^4X \times \text{Sh}(DX).$$

There are two ways to proceed. *Option 1* keeps this absolute indexing and accepts a target that contains a fourth-order level (section 8); the transformation of the first version omitted that level and was therefore ill-typed. *Option 2* fixes an outcome base B once and for all and lets the three levels live over B whatever object the comonad is evaluated at (section 9). Option 2 preserves exactly three hierarchical levels and is the main construction of this paper; it admits a strict mixed distributive law (theorem 9.8). Option 1 is developed to show precisely what goes wrong without a fixed base.

Contributions.

1. A base-relative three-level hierarchy over a fixed base B , with coherent contexts, the canonical hierarchy $\sigma(\delta_h) = (\rho_h, \delta_{\rho_h}, \delta_{\delta_{\rho_h}}, \delta_h)$ as an affine section, and a strict mixed distributive law satisfying all four Beck axioms, natural in the base (section 9).
2. A coherent lifting problem for claims against the tower, with level-wise obstruction certificates, strictness and monotonicity of failure levels, and the behavior of obstructions under pooling (section 10); and an obstruction tower with indeterminacy, in which a secondary obstruction, defined only once the primary one vanishes, depends on the choice of intermediate lift (section 10.4).
3. The coalgebras, lifted monad and bialgebras of the base-relative law, a precise statement of the ambient 2-category and of what two-dimensional monad theory does and does not add, and convex-order Beck comparisons for the absolute tower showing that the naive convex-order relaxation fails, through an orientation obstruction and a comparability obstruction (section 12).
4. A general *context comonad* $C_E(X) = X \times E(X)$ for any endofunctor E of **Set** (theorem 6.1), specialized to the three-level tower C_3 .
5. A hypothesis functor (Hyp, q) with q natural, so that the compatibility square is naturality (section 3); the context lift $J = \text{Hyp}(j)$ and the distribution-level lift K satisfy all coherence conditions by proof (propositions 6.7 and 7.2).

6. A normalized partial algebra $(\oplus, \otimes, \text{Cond})$ on structured posteriors with corrected laws (section 4).
7. An affine first isomorphism theorem for observational equivalence, and an exact characterization of when it is a congruence for $\otimes$ and for Cond (section 5); identifiability of source weights, a sharp supports-only interval, and robust decomposition under unmodeled contamination with a certified lower bound on the unmodeled weight (sections 5.2 and 5.3); and their sequential form for generated tokens: a likelihood projection onto declared contexts with bands, an unknown certificate, and a level-two test of blending against switching (section 5.4).
8. For the absolute tower, two mixed distributive laws λ^σ and $\bar{\lambda}$ with a complete account of which Beck axioms each satisfies (section 8, summarized in table 2), together with a structural lemma showing that the unit axiom forbids point-mass hierarchy at level one (lemma 8.1).
9. A proof that composite identifiability trivializes the hypothesis space, and a proof that the likelihood-factorization “iff” holds under nondegeneracy alone (remark 11.2 and theorem 11.3).

Structure of the argument. The categorical results form a progression:

- (1) a strict mixed distributive law for the absolute tower fails (section 8);
- (2) the base-relative reformulation restores a strict law satisfying all four Beck axioms (theorem 9.8), which remains the control result of the paper;
- (3) one may ask whether an order enrichment recovers a weaker law for the absolute tower; for the convex order the answer is no, because of an orientation obstruction and a comparability obstruction (propositions 12.8 to 12.10);
- (4) this leaves open which richer enrichment, comparing coherent lifts (section 10) and not only barycentric spreads, could do so (remark 12.12).

Relation to earlier versions. The first version claimed that a three-level transformation $\lambda^{[3]}$ satisfies the unit, counit and comultiplication axioms exactly and extends to a weak distributive law with a barycentric 2-cell. That transformation omitted the D^4X component of its target, and the claims do not hold; the corrected statements for the absolute tower are theorems 8.3, 8.6 and 8.7. The second version corrected the typing but left the existence of a full law open. Version 3 resolved the question for the base-relative tower; version 4 added the lifting problem; version 5 corrects the 2-categorical discussion and several empirical attributions. A full list of changes is in appendix B.

2 Preliminaries

2.1 The finite distribution monad

Let D be the finitely supported distribution monad on **Set**: $D(X)$ is the set of finitely supported probability measures on X , $D(f) = f_*$ is pushforward, $\eta_X(x) = \delta_x$, and $\mu_X: D^2X \rightarrow DX$ is barycenter, $\mu_X(\sum_i a_i \delta_{\rho_i}) = \sum_i a_i \rho_i$. Its Kleisli category restricted to finite sets is the category of finite stochastic matrices; see [13, 21] for the Markov-category perspective.

We work on **Set** rather than on finite sets: even for finite X with $|X| \geq 2$, the set $D(X)$ is an infinite simplex, so D does not restrict to an endofunctor of finite sets. Finiteness enters only through finite supports.

Each $D(Y)$ is a convex set, and every $D(f)$ is affine. A product of convex sets is convex componentwise, and a product of affine maps is affine.

2.2 Mixed distributive laws

Definition 2.1 (Mixed distributive law [10, 31]). Let (T, η, μ) be a monad and (C, ε, δ) a comonad on the same category. A *mixed distributive law* is a natural transformation $\lambda: TC \Rightarrow CT$ such that

$$\begin{array}{ll} \text{(U)} & \lambda \circ \eta_C = C\eta, \\ \text{(Co)} & \varepsilon_T \circ \lambda = T\varepsilon, \\ \text{(M)} & \lambda \circ \mu_C = C\mu \circ \lambda_T \circ T\lambda, \\ \text{(Cm)} & \delta_T \circ \lambda = C\lambda \circ \lambda_C \circ T\delta. \end{array}$$

Weak distributive laws between two monads, in the sense of Street and Böhm [36, 11, 12], keep the multiplication-type axioms and relax the unit-type ones; Garner [23] and Goy-Petrişan [24] give probabilistic and topological examples. We return to this comparison in [remark 8.9](#).

3 Structured hypotheses

Definition 3.1 (Hypothesis functor). A *hypothesis functor* is a functor $\text{Hyp}: \mathbf{Set} \rightarrow \mathbf{Set}$ with a natural transformation $q: \text{Hyp} \Rightarrow \mathbf{D}$, the *predictive map*. Naturality of q is the compatibility square

$$\begin{array}{ccc} \text{Hyp}(X) & \xrightarrow{q_X} & \mathbf{D}(X) \\ \text{Hyp}(f) \downarrow & & \downarrow f_* \\ \text{Hyp}(Y) & \xrightarrow{q_Y} & \mathbf{D}(Y) \end{array} \quad \text{i.e.} \quad q_Y \circ \text{Hyp}(f) = f_* \circ q_X.$$

Definition 3.2 (Structured hypotheses). Fix a set Θ of *structural descriptors*

$$\theta = (I_\theta, F_\theta, E_\theta, L_\theta, G_\theta, C_\theta) :$$

a sparse multi-index of directions, a factor graph, an evidence set, spline links, geometric data and conditional structure. Put

$$\text{Hyp}(X) = \Theta \times \mathbf{D}(X), \quad \text{Hyp}(f)(\theta, \rho) = (\theta, f_*\rho), \quad q_X(\theta, \rho) = \rho.$$

Lemma 3.3. *Hyp is a functor and q is natural.*

Proof. $\text{Hyp} = \Theta \times \mathbf{D}(-)$ is a product of a constant functor with $\mathbf{D}$, and q is the projection, which is natural because $q_Y(\text{Hyp}(f)(\theta, \rho)) = f_*\rho = f_*(q_X(\theta, \rho))$. $\square$

Remark 3.4 (Why descriptors are not transported). Directions v_j and spline links s_j are functions on X . Such data pull back along maps *into* X ; they cannot be pushed along an arbitrary $f: X \rightarrow Y$ (the expression $v_j \circ f$ is not defined), and conditional structure has no canonical pushforward along non-injective maps. In CausalIDView the descriptors parametrize a model on a fixed design, and the only datum that must move along outcome maps is the predictive distribution. [Definition 3.2](#) records exactly this. The *admissible* hypotheses $\text{Hyp}^{\text{adm}}(X) = \{(\theta, \rho_\theta)\}$, in which the predictive is the one produced by θ , form a subset that is generally not closed under $\text{Hyp}(f)$; all constructions below are carried out on the full functor Hyp .

Definition 3.5 (Shadow). $\text{Sh}(X) = \mathbf{D}(\text{Hyp}(X))$, the finite convex combinations $\Sigma = \sum_i \lambda_i \delta_{h_i}$ of hypotheses. Sh is a functor with $\text{Sh}(f) = \mathbf{D}(\text{Hyp}(f))$.

4 The structured posterior

Definition 4.1. $\text{Post}(X) = \text{D}(X \times \text{Hyp}(X))$. An element is $P = \sum_k \mu_k \delta_{(x_k, h_k)}$. The *outcome marginal* is $\pi_{X*} = \text{D}(\text{pr}_1): \text{Post}(X) \rightarrow \text{D}(X)$, and the *predictive* is $p_X = \mu_X \circ \text{D}(q_X \circ \text{pr}_2)$, so $p_X(P) = \sum_k \mu_k q_X(h_k)$.

For a nonzero finitely supported measure $m \geq 0$ write $\mathcal{N}(m) = m/\|m\|$ for its normalization. $\mathcal{N}$ absorbs positive scalars: $\mathcal{N}(cm) = \mathcal{N}(m)$ for $c > 0$.

4.1 Mixture

Definition 4.2. $P_1 \oplus_\lambda P_2 = \lambda P_1 + (1 - \lambda)P_2$ for $\lambda \in [0, 1]$.

Lemma 4.3 (Barycentric associativity). *For $\lambda, \mu \in [0, 1]$ with $\lambda\mu < 1$ and $\nu = \frac{(1-\lambda)\mu}{1-\lambda\mu}$,*

$$(P_1 \oplus_\lambda P_2) \oplus_\mu P_3 = P_1 \oplus_{\lambda\mu} (P_2 \oplus_\nu P_3).$$

If $\lambda = \mu = 1$ both sides equal P_1 for any choice of the inner weight.

Proof. The left side is $\lambda\mu P_1 + (1 - \lambda)\mu P_2 + (1 - \mu)P_3$. On the right, $(1 - \lambda\mu)\nu = (1 - \lambda)\mu$ and $(1 - \lambda\mu)(1 - \nu) = 1 - \lambda\mu - (1 - \lambda)\mu = 1 - \mu$. The boundary case is immediate. $\square$

4.2 Composition

Definition 4.4 (Compositional structure). A compositional structure on $\text{Hyp}(X)$ consists of a relation $\perp$ (compatibility), a map $\otimes$ defined on compatible pairs, and an element h_{id} , such that:

(A1) (ternary admissibility and associativity) $h_1 \perp h_2$ and $(h_1 \otimes h_2) \perp h_3$ if and only if $h_2 \perp h_3$ and $h_1 \perp (h_2 \otimes h_3)$, and then $(h_1 \otimes h_2) \otimes h_3 = h_1 \otimes (h_2 \otimes h_3)$;

(A2) (identity) $h_{\text{id}} \perp h$, $h \perp h_{\text{id}}$ and $h_{\text{id}} \otimes h = h \otimes h_{\text{id}} = h$ for all h .

Definition 4.5 (Composition of posteriors). The *unnormalized product* is

$$(P_1 \otimes P_2)(x, h) = \sum_{\substack{h_1 \perp h_2 \\ h_1 \otimes h_2 = h}} P_1(x, h_1) P_2(x, h_2),$$

a finitely supported measure of total mass $Z(P_1, P_2) \in [0, 1]$. If $Z(P_1, P_2) > 0$ put $P_1 \otimes P_2 = \mathcal{N}(P_1 \otimes P_2) \in \text{Post}(X)$; otherwise $P_1 \otimes P_2$ is undefined.

Remark 4.6. Without normalization the product is not a probability distribution: if $\perp$ is total its mass is $\sum_x m_1(x)m_2(x)$, where $m_i = \pi_{X*}(P_i)$, which is strictly less than 1 unless the marginals are equal point masses.

Lemma 4.7. $\otimes$ extends bilinearly to finitely supported measures and is associative. Consequently $(P_1 \otimes P_2) \otimes P_3 = P_1 \otimes (P_2 \otimes P_3)$, and either side is defined exactly when $P_1 \otimes P_2 \otimes P_3 \neq 0$.

Proof. Bilinearity is clear from the formula. Expanding both triple products gives sums over triples with $(h_1 \otimes h_2) \otimes h_3 = h$, respectively $h_1 \otimes (h_2 \otimes h_3) = h$, over the admissible triples; by (A1) these index sets and terms coincide. Since $\otimes$ is bilinear and $\mathcal{N}$ absorbs positive scalars, $\mathcal{N}(\mathcal{N}(m) \otimes m') = \mathcal{N}(m \otimes m')$, so both normalized products equal $\mathcal{N}(P_1 \otimes P_2 \otimes P_3)$. If $P_1 \otimes P_2 = 0$ the triple product is 0, which gives the definedness statement. $\square$

Proposition 4.8 (Distributivity). *Let $Z_{i3} = Z(P_i, P_3)$ and suppose $Z_{13}, Z_{23} > 0$. Then*

$$(P_1 \oplus_\lambda P_2) \otimes P_3 = (P_1 \otimes P_3) \oplus_{\lambda''} (P_2 \otimes P_3), \quad \lambda'' = \frac{\lambda Z_{13}}{\lambda Z_{13} + (1 - \lambda) Z_{23}}.$$

Proof. By bilinearity $(\lambda P_1 + (1 - \lambda) P_2) \otimes P_3 = \lambda Z_{13} (P_1 \otimes P_3) + (1 - \lambda) Z_{23} (P_2 \otimes P_3)$; normalize. $\square$

The weight is $\lambda'' = \lambda$ only when $Z_{13} = Z_{23}$ (or $\lambda \in \{0, 1\}$).

4.3 Conditioning

Definition 4.9. For evidence e with likelihood $L(e \mid \cdot) : \text{Hyp}(X) \rightarrow [0, \infty)$, put $Z_e(P) = \sum_{(x,h)} L(e \mid h) P(x, h)$ and, when $Z_e(P) > 0$, $\text{Cond}_e(P) = \mathcal{N}(L(e \mid \cdot) P)$.

Proposition 4.10 (Chain rule). *Suppose $L(e_1 \wedge e_2 \mid h) = L(e_1 \mid h) L(e_2 \mid h, e_1)$, and write $\text{Cond}_{e_2|e_1}$ for conditioning with likelihood $L(e_2 \mid \cdot, e_1)$. Then $\text{Cond}_{e_2|e_1}(\text{Cond}_{e_1} P) = \text{Cond}_{e_1 \wedge e_2}(P)$ whenever the right side is defined. In particular, if e_1, e_2 are conditionally independent given h , then $\text{Cond}_{e_2} \circ \text{Cond}_{e_1} = \text{Cond}_{e_1 \wedge e_2}$.*

Proof. $\mathcal{N}(L_2 \cdot \mathcal{N}(L_1 P)) = \mathcal{N}(L_2 L_1 P)$. $\square$

Proposition 4.11 (Conditioning and mixture). *If $Z_1 = Z_e(P_1) > 0$ and $Z_2 = Z_e(P_2) > 0$, then $\text{Cond}_e(P_1 \oplus_\lambda P_2) = \text{Cond}_e(P_1) \oplus_{\lambda'} \text{Cond}_e(P_2)$ with $\lambda' = \lambda Z_1 / (\lambda Z_1 + (1 - \lambda) Z_2)$.*

Proof. $L \cdot (\lambda P_1 + (1 - \lambda) P_2) = \lambda Z_1 \text{Cond}_e(P_1) + (1 - \lambda) Z_2 \text{Cond}_e(P_2)$; normalize. $\square$

Proposition 4.12 (Conditioning and composition). *If $L(e \mid h_1 \otimes h_2) = L(e \mid h_1) L(e \mid h_2)$ for all compatible pairs, then $\text{Cond}_e(P_1 \otimes P_2) = \text{Cond}_e(P_1) \otimes \text{Cond}_e(P_2)$, and the left side is defined exactly when the right side is.*

Proof. Pointwise, $L(e \mid h_1 \otimes h_2) P_1(x, h_1) P_2(x, h_2) = (LP_1)(x, h_1) (LP_2)(x, h_2)$, so $L \cdot (P_1 \otimes P_2) = (LP_1) \otimes (LP_2)$. Using that $\mathcal{N}$ absorbs positive scalars and $\otimes$ is bilinear, $\text{Cond}_e(P_1 \otimes P_2) = \mathcal{N}(L \cdot (P_1 \otimes P_2)) = \mathcal{N}((LP_1) \otimes (LP_2)) = \mathcal{N}(\text{Cond}_e P_1 \otimes \text{Cond}_e P_2)$. Both sides are defined iff $(LP_1) \otimes (LP_2) \neq 0$. $\square$

Remark 4.13 (Partial algebra). $\text{Post}(X)$ is a convex set with a partial, normalized, associative composition and a partial conditioning action. It is not a semiring: $\oplus$ distributes over $\otimes$ only with the reweighting of [proposition 4.8](#).

5 Observational equivalence

5.1 Equivalence and congruence

Definition 5.1. $P_1 \equiv P_2$ iff $\pi_{X*}(P_1) = \pi_{X*}(P_2)$.

Theorem 5.2 (Affine first isomorphism theorem). *π_{X*} is affine. The relation $\equiv$ is a congruence for convex combinations, so $\text{Post}(X)/\equiv$ carries a unique convex structure making the quotient map affine, and π_{X*} induces an affine injection $\text{Post}(X)/\equiv \hookrightarrow \text{D}(X)$. If $\text{Hyp}(X) \neq \emptyset$ it is a bijection, so $\text{Post}(X)/\equiv \cong \text{D}(X)$ as convex sets.*

Proof. $\pi_{X*} = \text{D}(\text{pr}_1)$ is affine. If $P_i \equiv P'_i$ then $\pi_{X*}(\sum a_i P_i) = \sum a_i \pi_{X*}(P_i) = \pi_{X*}(\sum a_i P'_i)$, so $\equiv$ is a congruence and the induced map is well defined, affine and injective. For surjectivity pick $h_0 \in \text{Hyp}(X)$; then $\pi_{X*}(\sum_x \rho(x) \delta_{(x, h_0)}) = \rho$. $\square$

Proposition 5.3 (Congruence for $\oplus$). *If $P_1 \equiv P'_1$ and $P_2 \equiv P'_2$ then $P_1 \oplus_\lambda P_2 \equiv P'_1 \oplus_\lambda P'_2$.*

Proposition 5.4 (Congruence for $\otimes$). *If $\perp$ is total, then $\pi_{X*}(P_1 \otimes P_2) = \mathcal{N}(m_1 m_2)$ with $m_i = \pi_{X*}(P_i)$, so $\equiv$ is a congruence for $\otimes$ (no condition on q is needed). If $\perp$ is not total this can fail.*

Proof. With $\perp$ total, $\sum_h (P_1 \otimes P_2)(x, h) = \sum_{h_1, h_2} P_1(x, h_1) P_2(x, h_2) = m_1(x) m_2(x)$. For the failure, take $X = \{x, y\}$, hypotheses h_1, h_2, h_3 with $h_1 \perp h_3$ but not $h_2 \perp h_3$, and $P_1 = \frac{1}{2}\delta_{(x, h_1)} + \frac{1}{2}\delta_{(y, h_2)}$, $P'_1 = \frac{1}{2}\delta_{(x, h_2)} + \frac{1}{2}\delta_{(y, h_1)}$, $P_2 = \frac{1}{2}\delta_{(x, h_3)} + \frac{1}{2}\delta_{(y, h_3)}$. Then $P_1 \equiv P'_1$ but the marginals of $P_1 \otimes P_2$ and $P'_1 \otimes P_2$ are δ_x and δ_y . $\square$

Proposition 5.5 (Congruence for Cond). *Let $|X| \geq 2$. Then $\equiv$ is a congruence for Cond_e (on all P where both sides are defined) if and only if $L(e \mid \cdot)$ is constant on $\text{Hyp}(X)$.*

Proof. If L is constant and positive, $\text{Cond}_e = \text{id}$. Conversely, suppose $L(e \mid h_1) = \ell_1 \neq \ell_2 = L(e \mid h_2)$, so $\ell_1 + \ell_2 > 0$. With distinct $x_1, x_2 \in X$ let $P = \frac{1}{2}\delta_{(x_1, h_1)} + \frac{1}{2}\delta_{(x_2, h_2)}$ and $P' = \frac{1}{2}\delta_{(x_1, h_2)} + \frac{1}{2}\delta_{(x_2, h_1)}$. Then $P \equiv P'$, but the marginals of $\text{Cond}_e P$ and $\text{Cond}_e P'$ are $(\ell_1, \ell_2)/(\ell_1 + \ell_2)$ and $(\ell_2, \ell_1)/(\ell_1 + \ell_2)$, which differ. $\square$

Remark 5.6. In particular, requiring $L(e \mid h) = \ell(e \mid q_X(h))$ does *not* make $\equiv$ a congruence for conditioning: in the counterexample $q_X(h_1) \neq q_X(h_2)$ is allowed. With $\ell_1 = 2, \ell_2 = 1$ the conditioned marginals are $(\frac{2}{3}, \frac{1}{3})$ and $(\frac{1}{3}, \frac{2}{3})$.

5.2 Identifiability of decompositions

Theorem 5.2 says that outcome data determine only the image in $\text{D}(X)$; the hypothesis that produced them is quotiented away. We now ask when the weights of known sources can nevertheless be recovered. In this subsection and the next, X is finite and $A_1, \dots, A_k \in \text{D}(X)$ are fixed *sources*, for instance the predictives $q_X(h_i)$ of hypotheses or the admissible posteriors S_1 of [section 10](#). For $x \in \text{D}(X)$ the set of *decompositions* of x is

$$W(x) = \left\{ \lambda \in \Delta_k : \sum_i \lambda_i A_i = x \right\},$$

a polytope, nonempty if and only if $x \in \text{conv}(A_1, \dots, A_k)$. The *exclusive region* of source i is $R_i = \{b \in X : A_i(b) > 0, A_j(b) = 0 \text{ for } j \neq i\}$.

Proposition 5.7 (Identifiability).

- (i) $W(x)$ has at most one point for every x if and only if $A_1, \dots, A_k$ are affinely independent. In that case $x \mapsto \lambda(x)$ is affine on $\text{conv}(A_1, \dots, A_k)$.
- (ii) For any sources, if $A_i(R_i) > 0$ then $\lambda_i = x(R_i)/A_i(R_i)$ for every $\lambda \in W(x)$.
- (iii) At most $|X|$ sources can be affinely independent.
- (iv) For each i , $\{\lambda_i : \lambda \in W(x)\}$ is a closed interval whose endpoints are the values of two linear programs over $W(x)$.
- (v) Given $\lambda \in W(x)$, an observation b with $x(b) > 0$ came from source i with probability $\lambda_i A_i(b)/x(b)$; as λ ranges over $W(x)$ this ranges over a closed interval with linear-programming endpoints, reducing to a point when $W(x)$ is a point.

Proof. (i) If $\sum_i \lambda_i A_i = \sum_i \lambda'_i A_i$ with $\lambda, \lambda' \in \Delta_k$, then $c = \lambda - \lambda'$ satisfies $\sum_i c_i A_i = 0$ and $\sum_i c_i = 0$, which forces $c = 0$ under affine independence. Conversely, an affine dependence $c \neq 0$ gives, at $x = \sum_i \frac{1}{k} A_i$, the two decompositions $\frac{1}{k} \mathbf{1} \pm tc$ for small $t > 0$. The inverse of an injective affine map is affine. (ii) On R_i only source i has mass, so $x(R_i) = \lambda_i A_i(R_i)$. (iii) Affinely independent points of $D(X)$, an affine space of dimension $|X| - 1$, number at most $|X|$. (iv) $W(x)$ is a polytope and $\lambda \mapsto \lambda_i$ is linear. (v) $x(b)$ is fixed, so the probability is linear in λ . $\square$

The data of [example 10.16\(b\)](#) give an instance of non-identifiability: for the sources δ_0, δ_1, u on $X = \{0, 1\}$, the set $W(u)$ is the segment joining the weight vectors $(0, 0, 1)$ and $(\frac{1}{2}, \frac{1}{2}, 0)$. Even when $W(x)$ is a point, attribution of a *single* observation in an overlap of supports is genuinely probabilistic, as in (v); what is exact is the attribution of the mixture as a whole, namely the measures $\lambda_i A_i$.

Proposition 5.8 (Supports only). *Suppose only the supports of two sources are known: $\text{supp } A \subseteq S_A$, $\text{supp } B \subseteq S_B$, with $\text{supp } x \subseteq S_A \cup S_B$. Put $R_A = S_A \setminus S_B$, $R_B = S_B \setminus S_A$ and $R_{AB} = S_A \cap S_B$. Then the set of λ for which there exist A, B with these supports and $x = \lambda A + (1 - \lambda)B$ is exactly*

$$[x(R_A), 1 - x(R_B)].$$

This interval is $[\text{Bel}(\{A\}), \text{Pl}(\{A\})]$ for the Dempster–Shafer mass function $m(\{A\}) = x(R_A)$, $m(\{B\}) = x(R_B)$, $m(\{A, B\}) = x(R_{AB})$ [16, 33]. More generally, for k sources with known supports, let $T(b)$ be the set of sources whose support contains b and $m(T) = x(\{b : T(b) = T\})$; the feasible weight vectors are exactly the core $\{\lambda \in \Delta_k : \sum_{i \in T} \lambda_i \geq \text{Bel}(T) \text{ for all } T\}$ of the belief function.

Proof. On R_A the source B vanishes, so $\lambda A = x$ there and $\lambda \geq x(R_A)$; symmetrically $1 - \lambda \geq x(R_B)$. Conversely, for λ in the interval put $a = \lambda - x(R_A) \geq 0$ and $b = 1 - \lambda - x(R_B) \geq 0$, so $a + b = x(R_{AB})$; let λA equal x on R_A and $\frac{a}{x(R_{AB})}x$ on R_{AB} , and $(1 - \lambda)B$ equal x on R_B and $\frac{b}{x(R_{AB})}x$ on R_{AB} (both zero on R_{AB} if $x(R_{AB}) = 0$). For k sources, a weight vector is feasible iff the mass $x(b)$ at each outcome can be allocated to the sources in $T(b)$ with totals λ_i . By Gale’s supply–demand theorem [22] this holds iff $\sum_{i \in S} \lambda_i \leq x(\{b : T(b) \cap S \neq \emptyset\}) = \text{Pl}(S)$ for every set S of sources, which, since $\sum_i \lambda_i = 1$, is the core condition applied to the complement of S . $\square$

The three regimes are summarized as follows.

Known about the sources	Weights	Attribution of overlap mass
full distributions, affinely independent	unique (proposition 5.7(i))	exact probability
full distributions, affinely dependent	polytope, LP bounds	interval
supports only	core of a belief function	belief–plausibility interval

5.3 Robust decomposition under unmodeled contamination

All three regimes assume that the declared sources are complete. To allow for sources outside the model we use the contamination model of Huber [26]:

$$x = \sum_i \lambda_i A_i + \mu U, \quad \lambda \geq 0, \mu \geq 0, \sum_i \lambda_i + \mu = 1, \quad U \in D(X) \text{ unknown.}$$

Lemma 5.9. *(λ, μ) is compatible with x for some U if and only if $\sum_i \lambda_i A_i \leq x$ pointwise and $\sum_i \lambda_i = 1 - \mu$. Then $U = (x - \sum_i \lambda_i A_i)/\mu$ when $\mu > 0$.*

Proof. $\mu U = x - \sum_i \lambda_i A_i$ must be a nonnegative measure of mass μ . If $\mu = 0$, a nonnegative measure of mass zero vanishes. $\square$

Theorem 5.10 (Unmodeled weight). *Let $\mu_{\min}(x) = 1 - \max\{\sum_i \lambda_i : \lambda \geq 0, \sum_i \lambda_i A_i \leq x\}$.*

- (i) *The unmodeled weights compatible with x form the interval $[\mu_{\min}(x), 1]$.*
- (ii) *(Certificate) $1 - \mu_{\min}(x) = \min\{\langle y, x \rangle : y \geq 0, \langle y, A_i \rangle \geq 1 \text{ for all } i\}$. Every such y certifies $\mu \geq 1 - \langle y, x \rangle$.*
- (iii) *$\mu_{\min}(x) = 0$ if and only if $x \in \text{conv}(A_1, \dots, A_k)$; the level-one test of [theorem 10.8](#) is the case of zero budget.*
- (iv) *(Budget) Under an assumed budget $\mu \leq \bar{\mu}$, x is explainable iff $\mu_{\min}(x) \leq \bar{\mu}$, and sharp bounds on each λ_i , on μ , and on the unmodeled mass $x(b) - \sum_i \lambda_i A_i(b)$ at each outcome are the values of linear programs over $\{\lambda \geq 0 : \sum_i \lambda_i A_i \leq x, \sum_i \lambda_i \geq 1 - \bar{\mu}\}$.*

Proof. (i) By [lemma 5.9](#), compatible μ are $1 - \sum_i \lambda_i$ over the feasible λ ; if λ is feasible so is $t\lambda$ for $t \in [0, 1]$. (ii) Linear programming duality; the primal is feasible at $\lambda = 0$ and the dual at a large constant y . (iii) If $\sum_i \lambda_i = 1$ and $\sum_i \lambda_i A_i \leq x$, both sides have mass one, so they are equal. (iv) Direct from [lemma 5.9](#). $\square$

Proposition 5.11 (The distance test is one-sided). *$d(x, \text{conv}(A_1, \dots, A_k)) \leq 2\mu_{\min}(x)$ in ℓ^1 , so $d > 2\bar{\mu}$ certifies that x is not explainable within budget $\bar{\mu}$. The converse fails: for one uniform source on three outcomes and $x = (\frac{1}{2}, \frac{1}{2}, 0)$, $d = \frac{2}{3}$ while $\mu_{\min}(x) = 1$, certified by $y = (0, 0, 3)$.*

Proof. If $x = s + \mu U$ with $s = \sum_i \lambda_i A_i$ of mass $1 - \mu$ and $\mu < 1$, put $a = s/(1 - \mu) \in \text{conv}(A)$; then $x - a = \mu(U - a)$ has norm at most 2μ (if $\mu = 1$ the bound $d \leq 2$ is trivial). In the example, x vanishes where the source has mass $\frac{1}{3}$, so no positive multiple of the source fits under x ; and $\langle y, A \rangle = 1, \langle y, x \rangle = 0$. $\square$

Proposition 5.12 (Counting is not identifiable). *If unmodeled sources may be arbitrary elements of $D(X)$, then x is compatible with m unmodeled sources of total weight ν if and only if it is compatible with one unmodeled source of weight ν .*

Proof. $\sum_j \nu_j U_j = \nu U$ with $U = \sum_j (\nu_j / \nu) U_j \in D(X)$, and conversely take $m - 1$ copies of weight zero. $\square$

Remark 5.13 (What is certified). From the data alone one obtains $\mu_{\min}(x)$: at least this fraction of x is not explained by the declared sources, with the explicit certificate y of [theorem 5.10\(ii\)](#). Every upper bound, and every rejection, depends on the assumed budget $\bar{\mu}$; with $\bar{\mu} = 1$ nothing is rejected. A certificate can show that the unmodeled weight exceeds a budget, but by [proposition 5.12](#) not that more than one unmodeled mechanism is present. Counting becomes meaningful only when unmodeled sources are restricted to a family that is not closed under mixing; for point masses, for instance, the least number of unmodeled sources is the least support size of $x - \sum_i \lambda_i A_i$ over compatible λ . In short, attribution is complete relative to a declared list of sources and a declared budget, and the certificates report when the data contradict them; no test can detect an unmodeled source whose contribution lies inside the declared convex hull.

5.4 Sequential attribution of generated tokens

Sections 5.2 and 5.3 decompose one distribution $x \in \mathcal{D}(X)$. A generator that writes text with two contexts available does not produce one distribution; it produces a sequence $x_1, \dots, x_n$ of tokens in a vocabulary V , and at step t each declared source has its own predictive given the prefix. We now give a projection of the sequence onto the declared sources, with a confidence band for the share of each context and a certified lower bound for the part that neither context explains.

Sources and the sequential contamination model. For each t let $A_t = p(\cdot \mid x_{<t}, \text{context } A)$ and $B_t = p(\cdot \mid x_{<t}, \text{context } B)$ be the generator's predictives with context A or context B in its window, and let U_t be a declared *reference* predictive, for example the same generator with neither context. These are computable from three forward passes. Write $a_t = A_t(x_t)$, $b_t = B_t(x_t)$, $u_t = U_t(x_t)$ and $P \in \mathbb{R}^{n \times 3}$ for the matrix with rows (a_t, b_t, u_t) . The sequential analogue of the contamination model of section 5.3 is

$$x_t \mid x_{<t} \sim w_A A_t + w_B B_t + w_U U_t, \quad w = (w_A, w_B, w_U) \in \Delta_3,$$

with constant weights. The reported coordinates are the *context share* $\theta = w_A/(w_A + w_B)$ and the *unexplained weight* $\mu = w_U$.

Definition 5.14 (Likelihood projection). The projection of $x_{1:n}$ onto the declared sources is

$$\pi(x_{1:n}) = \arg \max_{w \in \Delta_3} \ell(w), \quad \ell(w) = \sum_{t=1}^n \log(Pw)_t.$$

If the sources do not depend on the prefix ($A_t = A$, $B_t = B$, $U_t = U$), then $\ell(w)/n = -\text{KL}(\hat{x} \parallel w_A A + w_B B + w_U U) - H(\hat{x})$ for the empirical distribution $\hat{x}$, so π is the reverse information projection of $\hat{x}$ onto $\text{conv}(A, B, U)$ [15]. In particular $\pi(\hat{x}) \in W(\hat{x})$ whenever $\hat{x} \in \text{conv}(A, B, U)$, so π extends the decompositions of section 5.2 to data that do not lie exactly in the hull.

Proposition 5.15 (Sample identifiability).

- (i) ℓ is concave on Δ_3 , and the EM iteration $w_i \leftarrow \frac{1}{n} \sum_t w_i P_{ti} / (Pw)_t$ increases ℓ and converges to its maximum value [45, 43].
- (ii) If no $c \neq 0$ with $c_A + c_B + c_U = 0$ satisfies $Pc = 0$, then ℓ is strictly concave on Δ_3 and $\pi(x_{1:n})$ is a single point. If such a c exists, then ℓ is constant along the direction c , and an interior maximiser lies on a segment of maximisers.
- (iii) In particular, if $U_t = \alpha A_t + (1 - \alpha) B_t$ for all t and some fixed $\alpha \in [0, 1]$, then $c = (\alpha, 1 - \alpha, -1)$ is such a direction: the likelihood is flat in μ and no amount of data separates unexplained text from a blend of the contexts. This is the sequential form of proposition 5.7(i) and of the last sentence of the remark after proposition 5.12.

Proof. (i) $w \mapsto (Pw)_t$ is linear and $\log$ is concave; EM for mixture weights with fixed components is an alternating minimisation of a divergence, whose values converge to the minimum [43]. (ii) For w in the relative interior and c tangent to the simplex, $\frac{d^2}{ds^2} \ell(w + sc) = -\sum_t (Pc)_t^2 / (Pw)_t^2$, which is negative unless $Pc = 0$; if $Pc = 0$ then $\ell(w + sc) = \ell(w)$. (iii) $Pc = \alpha a + (1 - \alpha)b - u = 0$. $\square$

The next statement shows that near-dependence is already enough to lose the unexplained weight, and that it arises when the reference is folded into the context models instead of being kept as its own coordinate.

Proposition 5.16 (Absorbed reference). *Let A_t, B_t be the document parts of the two contexts and suppose each full context predictive mixes its document part with the reference: $A'_t = \beta A_t + (1 - \beta)U_t$ and $B'_t = \beta B_t + (1 - \beta)U_t$ with $\beta \in (0, 1)$. Suppose a fraction f of the tokens lie outside both documents ($a_t = b_t = 0 < u_t$) and that on the other tokens $a_t + b_t > 0 = u_t$. Then the unexplained weight estimated in the coordinates (A, B, U) is $\hat{\mu} = f$, while in the coordinates (A', B', U) it is*

$$\hat{\mu}' = \max\left(0, \frac{f - (1 - \beta)}{\beta}\right),$$

which is zero whenever $f \leq 1 - \beta$.

Proof. In the coordinates (A, B, U) , $\ell(w) = \sum_{a_t+b_t>0} \log(w_A a_t + w_B b_t) + \sum_{a_t=b_t=0} \log(w_U u_t)$. Writing $w_A + w_B = 1 - \mu$, the first sum is $(1 - f)n \log(1 - \mu)$ plus a function of θ alone, and the second is $fn \log \mu$ plus a constant, so $\hat{\mu} = f$. In the coordinates (A', B', U) the tokens outside both documents have likelihood $((1 - \beta)(1 - \mu) + \mu)u_t$, so up to terms in θ

$$\ell/n = (1 - f) \log(1 - \mu) + f \log(1 - \beta + \beta\mu),$$

whose derivative $-\frac{1-f}{1-\mu} + \frac{f\beta}{1-\beta+\beta\mu}$ vanishes at $\mu = (f - (1 - \beta))/\beta$ and is negative on $[0, 1)$ when $f \leq 1 - \beta$. $\square$

Proposition 5.16 is the design rule for the projection: the coordinates should be the parts of each source that the others do not already contain. With $\beta = \frac{1}{2}$, a generation that copies half of its tokens from a document outside the context window has $\hat{\mu}' = 0$; this is what the experiment below finds for the full context models.

Bands. For a scalar coordinate let $\ell_\theta(g) = \max\{\ell(w) : w_A/(w_A + w_B) = g\}$ and $\ell_\mu(m) = \max\{\ell(w) : w_U = m\}$ be the profile log-likelihoods and $\hat{\ell} = \max \ell$. The band for the context share is

$$C_\theta = \{g \in [0, 1] : 2(\hat{\ell} - \ell_\theta(g)) \leq \kappa_\theta \chi_{1,0.95}^2\},$$

and the certified lower bound for the unexplained weight is

$$\mu_{10} = \min\{m \in [0, 1] : 2(\hat{\ell} - \ell_\mu(m)) \leq \kappa_\mu \chi_{1,0.90}^2\}.$$

Two corrections to Wilks' calibration are needed. The tokens are dependent and the model is misspecified, so the threshold is scaled by a design effect $\kappa \geq 1$, the ratio of the variance of block sums of the score to its value under independence (a sandwich estimate in the sense of [48], with block length $\lceil \sqrt{n} \rceil$). And $\mu = 0$ lies on the boundary of the parameter space, where the likelihood-ratio statistic is asymptotically an equal mixture of χ_0^2 and χ_1^2 [41, 46], so the one-sided 95% bound uses $\chi_{1,0.90}^2 = 2.706$.

What is certified. The statement $\mu_{10} > 0$ rejects the hypothesis that the sequence is a blend of the two contexts, at level 5% asymptotically, and its validity requires only that the blend model be correct under the null; it does not require U to be the true source of the unexplained text. This is the sequential counterpart of [theorem 5.10\(iii\)](#), the test with zero budget. A positive budget $\bar{\mu}$ and the value of μ itself are relative to the declared reference direction U , as [theorem 5.10\(iv\)](#) and the remark after [proposition 5.12](#) require: with an arbitrary per-token contaminant ($U_t = \delta_{x_t}$) every sequence is explained by $\mu = 1$, which is the sequential form of the interval $[\mu_{\min}, 1]$. In practice the budget is calibrated: $\bar{\mu}$ is the 95% quantile of μ_{10} on generations known to use only the two contexts, so that paraphrase, which also lies in the reference direction, is not reported as unexplained.

Blend versus switching is a level-two distinction. A generator can use both contexts in two different ways. It can *blend* them, drawing each token from one mixture $\lambda A_t + (1 - \lambda)B_t$, or it can *switch* between them in spans, drawing x_t from A_t or B_t according to a hidden two-state chain z_t . If the chain is stationary with $\Pr(z_t = A) = \lambda$, each token is drawn from A_t with marginal probability λ in both cases, and the constant-weight projection π targets the same $\theta = \lambda$. They differ in the second-order state: a blend has a single posterior weight, $\Pi = \delta_\lambda$, while switching carries a distribution over weights concentrated near 0 and 1. This is [example 9.11](#) for generated text: pooling (here, the constant-weight projection) returns the same level-one state and hides the difference, which survives at level two. The test uses the two-state hidden Markov model with free transition probabilities p_{AB}, p_{BA} , whose memoryless case $p_{AB} = 1 - \lambda, p_{BA} = \lambda$ is exactly the blend, and the statistic

$$S = 2 \left(\max_{p_{AB}, p_{BA}} \log L_{\text{HMM}} - \max_{\lambda} \log L_{\text{blend}} \right) \geq 0.$$

Under the null the transition parameters are not identified, so the null distribution is not χ^2 [44]; the threshold is the 95% quantile of S over blends simulated from the declared sources.

Per-token attribution. Given $\hat{w}$, token t came from source i with probability $r_t(i) = \hat{w}_i P_{ti} / (P\hat{w})_t$, as in [proposition 5.7\(v\)](#); under switching, forward-backward smoothing in the fitted chain replaces $\hat{w}$ by the posterior of z_t . As in [section 5.2](#), what is exact is the attribution of the whole sequence; a single token in the overlap of two sources is attributed only probabilistically.

Decision rule. With budget $\bar{\mu}$ and switching threshold s^* , report

condition (checked in order)	report
$\mu_{\text{lo}} > \bar{\mu}$	<i>unknown</i> : a certified share $\geq \mu_{\text{lo}}$ is explained by neither context
$S > s^*$	<i>switching</i> between A and B , with θ and C_θ as time shares
$\min C_\theta \geq 0.9$	<i>from A</i>
$\max C_\theta \leq 0.1$	<i>from B</i>
otherwise	<i>blend</i> with share $\theta \in C_\theta$

Every report carries the band C_θ and the interval $[\mu_{\text{lo}}, 1]$ for the unexplained weight; the decision is a summary of them, not a replacement.

Experiment. We tested the projection in a retrieval-style setting on real news text: contexts A and B are two articles from the BASIL corpus [40], either about different events (“easy”) or two outlets’ articles about the same event (“hard”). The reference U is a bigram model trained on the remaining articles and 40,000 neutral headlines [39]; each context predictive is $\frac{1}{2}(\text{document bigram}) + \frac{1}{2}U$, and the coordinates are the document parts and U , as [proposition 5.16](#) requires. Generated text with known truth: a span copied from A or from B with 20% of the tokens replaced by samples from U (paraphrase noise); switching between spans of 20–50 tokens; blends sampled token by token at $\lambda = 0.2, 0.5, 0.8$; a span from a third article not in the context (*unknown*); and half A , half third article. Forty generations per case and pair type; the budget and threshold were calibrated on separate generations. [Table 1](#) gives the fraction of correct reports.

Text from a third article is reported as *unknown* in every case, with mean certified share $\mu_{\text{lo}} = 0.78\text{--}0.80$ at $n = 100$. The 95% band for the context share covers the true blend weight in 95–100% of generations at $n = 200$, with width 0.17–0.24. At $n = 200$ the hard pairs, two outlets writing about the same event, are attributed as well as the easy ones (copies 1.00/0.95, switching

	copy A	copy B	switch	blend	unk.	half unk.	C_θ cover	C_θ width	false unk.
$n = 50$, easy	0.78	0.78	0.85	0.88	1.00	0.90	0.85–1.00	0.33–0.40	0.05
$n = 50$, hard	0.50	0.53	0.70	0.87	1.00	1.00	0.93–0.97	0.37–0.45	0.09
$n = 100$, easy	0.93	0.97	0.90	0.91	1.00	1.00	0.88–0.97	0.23–0.28	0.07
$n = 100$, hard	0.85	0.82	0.93	0.86	1.00	1.00	0.95–0.97	0.26–0.33	0.05
$n = 200$, easy	0.97	0.93	0.95	0.92	1.00	1.00	0.97–1.00	0.17–0.20	0.03
$n = 200$, hard	1.00	0.95	0.95	0.95	1.00	1.00	0.95–1.00	0.19–0.24	0.01
$n = 100$, easy, <i>full</i>	0.97	0.97	1.00	0.97	1.00	0.00	0.90–1.00	0.24–0.31	0.00
$n = 100$, hard, <i>full</i>	0.88	0.88	1.00	0.89	1.00	0.03	0.93–0.97	0.28–0.37	0.00

Table 1: Context attribution of generated text (fraction of 40 generations per case reported correctly; “blend” is the mean over $\lambda = 0.2, 0.5, 0.8$; coverage and width of the 95% band C_θ over the three blends; “false unk.” is the rate of *unknown* reports on text that uses only the two contexts). Rows marked *full* use the full context models as coordinates, which absorbs the reference (proposition 5.16).

0.95). With 50 tokens the same-event pairs are hard: only about half of the copies are reported as pure copies, and most of the rest are reported as blends, a conservative error: in 27 of these 29 cases the band still contains the true share. The false-unknown rate on known-context text is 0.01–0.07 at $n \geq 100$, close to the 5% that the calibration of $\bar{\mu}$ allows. With the full context models as coordinates (rows *full*), the copies, blends and switching are attributed at least as well and nothing is falsely reported unknown, but text that is half from a third article is reported unknown in at most 2.5% of cases ($\hat{\mu}' = 0.08$ against 0.53 in the document coordinates), as proposition 5.16 predicts with $\beta = \frac{1}{2}$ and $f \approx \frac{1}{2}$.

Scope. The generations in this experiment come from bigram context models and from copied spans of real articles, not from a large language model, so the experiment tests the statistics of the projection (identifiability, band coverage, the unknown certificate and the level-two test) and not the behaviour of a particular generator. For a language model the procedure is the same, with A_t , B_t , U_t read from three forward passes. The bands assume constant weights within the sequence; for long generations they should be computed on windows, and the switching test then becomes a test for change points.

6 The three-level tower and the context comonad

6.1 The three posterior levels

For a set X let $H_3(X) = D(X) \times D^2(X) \times D^3(X)$, with elements (ρ, Π, Ω) . The barycentric projections $b_X = \mu_X: D^2X \rightarrow DX$ and $b_X^{(2)} = \mu_{DX}: D^3X \rightarrow D^2X$ form the tower $D^3X \rightarrow D^2X \rightarrow DX$. By monad associativity $\mu_X \circ \mu_{DX} = \mu_X \circ D(\mu_X)$, so the full collapse $D^3X \rightarrow DX$ does not depend on which projection is used at the top. The three levels are distinct state; the projections are observation maps.

6.2 A general context comonad

For a set X , an element c of a set A , and $y \in X$, write $j_{X,c}(y) = (y, c)$.

Theorem 6.1 (Context comonad). *Let $E: \mathbf{Set} \rightarrow \mathbf{Set}$ be any functor. Put $C_E(X) = X \times E(X)$ and $C_E(f) = f \times E(f)$, and define*

$$\varepsilon_X(x, c) = x, \quad \delta_X(x, c) = ((x, c), \kappa_X(c)), \quad \kappa_X(c) = E(j_{X,c})(c) \in E(C_EX),$$

where $j_{X,c}: X \rightarrow C_EX$. Then $(C_E, \varepsilon, \delta)$ is a comonad on $\mathbf{Set}$.

Proof. Write $C = C_E$. Functoriality is that of $\text{id} \times E$, and ε is natural because it is a product projection.

Naturality of δ . For $f: X \rightarrow Y$ we need $\delta_Y(fx, Ef c) = CCf(\delta_X(x, c))$. The second component of the right side is $E(Cf)(E(j_{X,c})(c)) = E(Cf \circ j_{X,c})(c)$. Now $(Cf \circ j_{X,c})(y) = (fy, Ef c) = (j_{Y, Ef c} \circ f)(y)$, so this equals $E(j_{Y, Ef c})(Ef c) = \kappa_Y(Ef c)$. First components agree trivially.

Counit laws. $\varepsilon_{CX} \delta_X(x, c) = (x, c)$. Also $C(\varepsilon_X) \delta_X(x, c) = (\varepsilon_X(x, c), E(\varepsilon_X)E(j_{X,c})(c)) = (x, c)$ since $\varepsilon_X \circ j_{X,c} = \text{id}_X$.

Coassociativity. Let $w = (x, c)$ and $c' = \kappa_X(c)$. We first note

$$\delta_X \circ j_{X,c} = j_{CX,c'} \circ j_{X,c}, \tag{1}$$

since both send y to $((y, c), c')$; this uses that $\kappa_X(c)$ depends only on c . Now $\delta_{CX} \delta_X(x, c) = ((w, c'), E(j_{CX,c'})(c')) = ((w, c'), E(j_{CX,c'} \circ j_{X,c})(c))$, while $C(\delta_X) \delta_X(x, c) = (\delta_X(w), E(\delta_X)(c')) = ((w, c'), E(\delta_X \circ j_{X,c})(c))$. These agree by (1). $\square$

Remark 6.2. If E is constant at a set A , then C_E is the reader (product) comonad $X \times A$. For general E the context is transported into C_EX by κ_X rather than copied.

6.3 The three-level context

Definition 6.3. $E_3(X) = D(X) \times D^2(X) \times D^3(X) \times \text{Sh}(X)$ and $C_3 = C_{E_3}$, so $C_3(X) = X \times DX \times D^2X \times D^3X \times \text{Sh}X$ with points $(x, \rho, \Pi, \Omega, \Sigma)$.

Corollary 6.4. C_3 is a comonad. Explicitly, for $c = (\rho, \Pi, \Omega, \Sigma)$,

$$\kappa_X(c) = ((j_{X,c})_* \rho, D^2(j_{X,c})(\Pi), D^3(j_{X,c})(\Omega), D(\text{Hyp}(j_{X,c}))(\Sigma)).$$

Each component of $E_3(Y)$ is of the form $D(Z)$, so $E_3(Y)$ is a convex set, and for every f the map $E_3(f)$ is affine.

Remark 6.5 (The absolute tower grows under D). Because the levels of C_3 are indexed by the object, $C_3(DX) = DX \times D(DX) \times D^2(DX) \times D^3(DX) \times \text{Sh}(DX) = DX \times D^2X \times D^3X \times D^4X \times \text{Sh}(DX)$. A transformation $DC_3X \rightarrow C_3DX$ must therefore supply a fourth-order component, and applying D again raises the order once more. In the notation of the first version, the components $P_M \in DX$, $\Pi_M \in D^2X$ and $\Omega_M \in D^3X$ fill only the first three factors. The transport in [corollary 6.4](#) is $\rho \mapsto (j_{X,c})_* \rho$, $\Pi \mapsto D^2(j_{X,c})\Pi$, $\Omega \mapsto D^3(j_{X,c})\Omega$; it is this re-indexing of the context over C_3X that makes the order grow.

6.4 The context lift

Definition 6.6. $J_{X,c} = \text{Hyp}(j_{X,c}): \text{Hyp}(X) \rightarrow \text{Hyp}(C_3X)$; explicitly $J_{X,c}(\theta, \rho) = (\theta, (j_{X,c})_* \rho)$.

Proposition 6.7 (Coherence of J). (CL1) $q_{C_3X} \circ J_{X,c} = (j_{X,c})_* \circ q_X$;

(CL2) $\text{Hyp}(\varepsilon_X) \circ J_{X,c} = \text{id}_{\text{Hyp}(X)}$;

(CL3) $\text{Hyp}(\delta_X) \circ J_{X,c} = J_{\mathbf{C}_3 X, \kappa_X(c)} \circ J_{X,c}$.
The shadow component of $\kappa_X(c)$ is $(J_{X,c})_ \Sigma$.*

Proof. (CL1) is naturality of q . (CL2) is functoriality of Hyp and $\varepsilon_X \circ j_{X,c} = \text{id}$. (CL3) is functoriality of Hyp applied to (1). $\square$

7 The distribution-level lift

For $k \geq 1$ let $\eta_X^{(k)} = \eta_{\mathbf{D}^k X} \circ \cdots \circ \eta_{\mathbf{D} X} : \mathbf{D} X \rightarrow \mathbf{D}^{k+1} X$, so $\eta_X^{(3)}(\rho) = \delta_{\delta_\rho}$.

Definition 7.1. $K_X^{[k]} : \text{Hyp}(X) \rightarrow \text{Hyp}(\mathbf{D}^k X)$, $K_X^{[k]}(\theta, \rho) = (\theta, \eta_X^{(k)}(\rho))$. Then $q_{\mathbf{D}^k X}(K_X^{[k]}h) = \eta_X^{(k)}(q_X h)$; for $k = 3$ this is $\delta_{\delta_{\rho_h}}$.

Proposition 7.2 (Naturality and K/J coherence). $K^{[k]} : \text{Hyp} \Rightarrow \text{Hyp} \circ \mathbf{D}^k$ is natural. In particular

$$\text{Hyp}(\mathbf{D}^3 j_{X,c}) \circ K_X^{[3]} = K_{\mathbf{C}_3 X}^{[3]} \circ J_{X,c}.$$

Proof. $\eta^{(k)}$ is a composite of whiskered units, hence natural: $\mathbf{D}^{k+1}(f) \circ \eta_X^{(k)} = \eta_Y^{(k)} \circ \mathbf{D}(f)$. Apply with $f = j_{X,c}$. $\square$

Remark 7.3. The coherence must use $\mathbf{D}^3 j_{X,c}$: $K_X^{[3]}$ lands in $\text{Hyp}(\mathbf{D}^3 X)$, on which $\text{Hyp}(\mathbf{D}^2 j)$ does not act. Note also that $K_X^{[3]}(h)$ is a point mass at the single point $\delta_{\delta_{\rho_h}}$ of $\mathbf{D}^3 X$; it is not the pushforward of ρ_h along $x \mapsto \delta_{\delta_x}$, which would be $\text{Hyp}(\eta\eta\eta)$.

8 Option 1: mixed distributive laws for the absolute tower

Throughout, $M = \sum_i a_i \delta_{(x_i, c_i)} \in \mathbf{DC}_3 X$ with $c_i = (\rho_i, \Pi_i, \Omega_i, \Sigma_i)$. The target is $\mathbf{C}_3 \mathbf{D} X = \mathbf{D} X \times \mathbf{D}(\mathbf{D} X) \times \mathbf{D}^2(\mathbf{D} X) \times \mathbf{D}^3(\mathbf{D} X) \times \text{Sh}(\mathbf{D} X)$.

8.1 What the unit axiom forces

Lemma 8.1. *If $\lambda : \mathbf{DC}_3 \Rightarrow \mathbf{C}_3 \mathbf{D}$ satisfies (U), then $\lambda(\delta_{(x,c)}) = (\delta_x, \mathbf{E}_3(\eta_X)(c))$. In particular its level-one component is $\mathbf{D}(\eta_X)(\rho) = \sum_y \rho(y) \delta_y$, which differs from δ_ρ whenever ρ is not a point mass. The two agree after barycentric collapse: $\mu_X(\mathbf{D}(\eta_X)\rho) = \rho = \mu_X(\delta_\rho)$.*

Proof. (U) states $\lambda(\eta_{\mathbf{C}_3 X}(x, c)) = \mathbf{C}_3(\eta_X)(x, c) = (\eta_X x, \mathbf{E}_3(\eta_X)c)$. $\square$

So a law that keeps a non-degenerate posterior ρ as a single hypothesis δ_ρ at level one cannot satisfy (U); the unit axiom requires level-one information to be spread over point hypotheses. This is the precise form of the tension between hierarchy and barycentric collapse.

8.2 The shift law

Definition 8.2 (Shift law). With $K_X^{[1]}(\theta, \rho) = (\theta, \delta_\rho)$,

$$\lambda_X^\sigma(M) = \left(\sum_i a_i \delta_{x_i}, \sum_i a_i \delta_{\rho_i}, \sum_i a_i \delta_{\Pi_i}, \sum_i a_i \delta_{\Omega_i}, \sum_i a_i (K_X^{[1]})_* \Sigma_i \right).$$

The components lie in $\mathbf{D} X$, $\mathbf{D}(\mathbf{D} X)$, $\mathbf{D}(\mathbf{D}^2 X) = \mathbf{D}^2(\mathbf{D} X)$, $\mathbf{D}(\mathbf{D}^3 X) = \mathbf{D}^3(\mathbf{D} X)$ and $\mathbf{D}(\text{Hyp} \mathbf{D} X) = \text{Sh}(\mathbf{D} X)$ respectively, so $\lambda_X^\sigma : \mathbf{DC}_3 X \rightarrow \mathbf{C}_3 \mathbf{D} X$.

Theorem 8.3. λ^σ is natural and satisfies (Co). It violates (U), (M) and (Cm). The defects in (U) and (M) at level one vanish after applying μ_X .

Proof. Naturality. Each component of λ^σ is the pushforward of M along a natural projection (followed, for the shadow, by the natural map $K^{[1]}$ and barycenter, both natural).

(Co). $\varepsilon_{DX}(\lambda^\sigma M) = \sum_i a_i \delta_{x_i} = D(\varepsilon_X)(M)$.

(U) fails. By lemma 8.1: $\lambda^\sigma(\delta_{(x,c)})$ has level one δ_ρ , while $C_3(\eta_X)(x, c)$ has $D(\eta_X)\rho$. For $X = \{0, 1\}$ and $\rho = \frac{1}{2}\delta_0 + \frac{1}{2}\delta_1$ these are δ_ρ and $\frac{1}{2}\delta_{\delta_0} + \frac{1}{2}\delta_{\delta_1}$.

(M) fails. Let $M' = \sum_\ell \nu_\ell \delta_{M_\ell}$ with $M_\ell = \sum_k a_{\ell k} \delta_{(x_{\ell k}, c_{\ell k})}$. The level-one component of $\lambda^\sigma \mu(M')$ is $\sum_{\ell, k} \nu_\ell a_{\ell k} \delta_{\rho_{\ell k}}$. On the other side, λ_{DX}^σ reads the level-one components $\sum_k a_{\ell k} \delta_{\rho_{\ell k}}$ of $\lambda^\sigma(M_\ell)$ and $C_3(\mu_X)$ applies $D(\mu_X)$, giving $\sum_\ell \nu_\ell \delta_{\sum_k a_{\ell k} \rho_{\ell k}}$. These differ whenever the $\rho_{\ell k}$ for fixed ℓ are not all equal, and both have barycenter $\sum_{\ell, k} \nu_\ell a_{\ell k} \rho_{\ell k}$.

(Cm) fails. Take $M = \frac{1}{2}\delta_{(x, c_1)} + \frac{1}{2}\delta_{(x, c_2)}$ with $\rho_1 \neq \rho_2$. The level-one component of $\delta_{DX}(\lambda^\sigma M)$ is $D(j_{DX, e})(\sum_i \frac{1}{2}\delta_{\rho_i}) = \sum_i \frac{1}{2}\delta_{(\rho_i, e)}$, where e is the pooled context of $\lambda^\sigma M$. On the other side, $C_3(\lambda^\sigma) \circ \lambda_{C_3 X}^\sigma \circ D(\delta_X)$ gives $\sum_i \frac{1}{2}\delta_{\lambda^\sigma((j_{X, c_i})_* \rho_i)} = \sum_i \frac{1}{2}\delta_{(\rho_i, e_i)}$ with $e_i = (\delta_{\rho_i}, \delta_{\Pi_i}, \delta_{\Omega_i}, (K^{[1]})_* \Sigma_i)$, and $e_i \neq e$. $\square$

Remark 8.4. The transformation $\lambda^{[3]}$ of the earlier version is λ^σ with the level-three component omitted and a shadow component of the wrong type. Its claimed unit and comultiplication theorems fail by theorem 8.3.

8.3 The averaging law

Let $\beta: D(E_3 Y) \rightarrow E_3 Y$ be the barycenter, computed componentwise by μ on each factor $D(Z)$; every $E_3(f)$ commutes with β because it is affine.

Definition 8.5 (Averaging law).

$$\bar{\lambda}_X(M) = \left(\sum_i a_i \delta_{x_i}, \sum_i a_i E_3(\eta_X)(c_i) \right),$$

i.e. $\bar{\lambda}_X = \langle D(\varepsilon_X), \beta \circ D(E_3(\eta_X) \circ \text{pr}_2) \rangle$.

Theorem 8.6. $\bar{\lambda}$ is natural and satisfies (U), (Co) and (M) exactly.

Proof. Naturality. For $f: X \rightarrow Y$, using that $E_3(Df)$ is affine and $Df \circ \eta_X = \eta_Y \circ f$,

$$E_3(Df) \left(\sum_i a_i E_3(\eta_X) c_i \right) = \sum_i a_i E_3(\eta_Y) (E_3(f) c_i),$$

which is the context component of $\bar{\lambda}_Y(D(C_3 f)M)$; first components agree by naturality of $D(\varepsilon)$.

(U). $\bar{\lambda}(\delta_{(x,c)}) = (\delta_x, E_3(\eta_X)c) = C_3(\eta_X)(x, c)$.

(Co). $\varepsilon_{DX} \circ \bar{\lambda}_X = D(\varepsilon_X)$ by definition.

(M). Let $M' = \sum_\ell \nu_\ell \delta_{M_\ell}$ with $M_\ell = \sum_k a_{\ell k} \delta_{(x_{\ell k}, c_{\ell k})}$, and put

$$P_\ell = \sum_k a_{\ell k} \delta_{x_{\ell k}}, \quad e_\ell = \sum_k a_{\ell k} E_3(\eta_X) c_{\ell k}.$$

The left side is $(\sum_\ell \nu_\ell P_\ell, \sum_\ell \nu_\ell e_\ell)$. On the right, $D(\bar{\lambda})M' = \sum_\ell \nu_\ell \delta_{(P_\ell, e_\ell)}$; then $\bar{\lambda}_{DX}$ gives

$$\left(\sum_\ell \nu_\ell \delta_{P_\ell}, \sum_\ell \nu_\ell E_3(\eta_{DX}) e_\ell \right),$$

	natural	(U)	(Co)	(M)	(Cm)	keeps δ_ρ at level one
shift law λ^σ	yes	no	yes	no	no	yes
averaging law $\bar{\lambda}$	yes	yes	yes	yes	no	no

Table 2: Status of the two mixed distributive laws $\mathbf{DC}_3 \Rightarrow \mathbf{C}_3\mathbf{D}$ (theorems 8.3, 8.6 and 8.7). By lemma 8.1, no law satisfying (U) keeps δ_ρ at level one.

and $\mathbf{C}_3(\mu_X)$ turns this into

$$\left(\sum_{\ell} \nu_{\ell} P_{\ell}, \sum_{\ell} \nu_{\ell} \mathbf{E}_3(\mu_X \circ \eta_{\mathbf{D}X}) e_{\ell} \right) = \left(\sum_{\ell} \nu_{\ell} P_{\ell}, \sum_{\ell} \nu_{\ell} e_{\ell} \right),$$

using that $\mathbf{E}_3(\mu_X)$ is affine and $\mu_X \circ \eta_{\mathbf{D}X} = \text{id}$. $\square$

Theorem 8.7. $\bar{\lambda}$ violates (Cm).

Proof. Let $X = \{0, 1\}$, fix Π, Ω, Σ , and let $c_i = (\delta_i, \Pi, \Omega, \Sigma)$ for $i = 0, 1$, $M = \frac{1}{2}\delta_{(0, c_0)} + \frac{1}{2}\delta_{(0, c_1)}$. Write $e_i = \mathbf{E}_3(\eta_X)c_i$ and $e = \frac{1}{2}e_0 + \frac{1}{2}e_1$, so $\bar{\lambda}M = (\delta_0, e)$ and the level-one component of e is $\frac{1}{2}\delta_{\delta_0} + \frac{1}{2}\delta_{\delta_1}$. The level-one component of $\delta_{\mathbf{D}X}(\bar{\lambda}M)$ is $\mathbf{D}(j_{\mathbf{D}X, e})$ of that, namely $\frac{1}{2}\delta_{(\delta_0, e)} + \frac{1}{2}\delta_{(\delta_1, e)}$.

On the other side, $\mathbf{D}(\delta_X)M = \sum_i \frac{1}{2}\delta_{((0, c_i), \kappa_X(c_i))}$. Since $\mathbf{E}_3(\bar{\lambda}_X)$ is affine and $\bar{\lambda}_X \circ \eta_{\mathbf{C}_3X} \circ j_{X, c_i} = j_{\mathbf{D}X, e_i} \circ \eta_X$ (by (U)), the context component of $\mathbf{C}_3(\bar{\lambda})\bar{\lambda}_{\mathbf{C}_3X}\mathbf{D}(\delta_X)(M)$ is $\sum_i \frac{1}{2}\mathbf{E}_3(j_{\mathbf{D}X, e_i})(e_i)$, whose level-one component is $\frac{1}{2}\delta_{(\delta_0, e_0)} + \frac{1}{2}\delta_{(\delta_1, e_1)}$. Since $e_0 \neq e \neq e_1$, the two sides differ. $\square$

Remark 8.8. In general the context component of the right side of (Cm) is $\sum_i a_i \mathbf{E}_3(j_{\mathbf{D}X, e_i})(e_i)$ and that of the left side is $\mathbf{E}_3(j_{\mathbf{D}X, e})(e)$ with $e = \sum_i a_i e_i$: comultiplication copies a *single* pooled context into $\mathbf{C}_3\mathbf{D}X$, whereas the right side keeps one context per sample. Both sides do agree after $\mathbf{C}_3(\varepsilon_{\mathbf{D}X})$ and after $\varepsilon_{\mathbf{C}_3\mathbf{D}X}$, but this follows from (Co) and naturality alone and carries no extra information.

Remark 8.9 (Relation to weak distributive laws). The weak distributive laws of Street and Böhm [36, 11, 12], used by Garner [23] and Goy–Petrişan [24], are laws between two monads that keep the multiplication axioms and weaken the unit axioms; the composite is obtained by splitting an idempotent. Neither λ^σ nor $\bar{\lambda}$ falls under that notion as stated: λ^σ fails a multiplication-type axiom (M), and $\bar{\lambda}$ fails the comonad analogue (Cm). A formulation in which a failing axiom is witnessed by a 2-cell would need a 2-categorical structure on the base (for instance an order or convex enrichment of the relevant hom-sets); on **Set** itself the only 2-cells between natural transformations are identities, so no non-identity 2-cell α is available. The earlier version’s claim of a weak distributive law with a barycentric 2-cell is therefore withdrawn. An Section 12.6 shows that the natural candidate order, the convex order, does not supply such 2-cells.

9 Option 2: the base-relative three-level hierarchy

By remark 6.5, the absolute tower cannot hold exactly three levels once the monad is applied. We now fix an outcome base B (in the application, the finite outcome space of CausalIDView) and let the hierarchy live over B regardless of the object at which the comonad is evaluated.

9.1 Coherent contexts

Definition 9.1 (Context algebra). $A_B = \mathbf{D}B \times \mathbf{D}^2B \times \mathbf{D}^3B \times \text{Sh}(B)$, with points $c = (\rho, \Pi, \Omega, \Sigma)$. It is a convex set componentwise, and the componentwise barycenter $\beta_B: \mathbf{D}(A_B) \rightarrow A_B$ (given by

μ on each factor DZ) makes (A_B, β_B) an algebra for D , namely the product of the free algebras (DZ, μ_Z) .

Definition 9.2 (Coherence). A context $c = (\rho, \Pi, \Omega, \Sigma) \in A_B$ is *coherent* if

$$\mu_B(\Pi) = \rho, \quad \mu_{DB}(\Omega) = \Pi, \quad D(q_B)(\Sigma) = \Pi.$$

Write $A_B^{\text{coh}} \subseteq A_B$ for the coherent contexts.

Coherence says that the barycentric observation maps of the tower $D^3B \rightarrow D^2B \rightarrow DB$ and the predictive of the shadow agree. The top level determines the lower two, but all three remain state.

Lemma 9.3. A_B^{coh} is closed under β_B ; hence it is a sub-algebra of (A_B, β_B) .

Proof. Each condition is an equation between affine maps ($\mu_B, \mu_{DB}, D(q_B)$ and the projections), and the solution set of such an equation is closed under convex combination. $\square$

9.2 The canonical hierarchy

Definition 9.4. $\sigma_B: \text{Sh}(B) \rightarrow A_B$,

$$\sigma_B(\Sigma) = (\mu_B(Dq_B \Sigma), Dq_B \Sigma, D(\eta_{DB})(Dq_B \Sigma), \Sigma).$$

On a single hypothesis, $\sigma_B(\delta_h) = (\rho_h, \delta_{\rho_h}, \delta_{\delta_{\rho_h}}, \delta_h)$, i.e. $\Pi_h = \delta_{\rho_h}$ and $\Omega_h = \delta_{\Pi_h}$.

Proposition 9.5. σ_B is affine, injective, lands in A_B^{coh} , and is natural in B : $A_f \circ \sigma_B = \sigma_{B'} \circ \text{Sh}(f)$ for $f: B \rightarrow B'$, where $A_f = Df \times D^2f \times D^3f \times \text{Sh}f$.

Proof. Each component is a composite of affine maps, and the last component is Σ , so σ_B is injective. Coherence: the first condition holds by definition, the second is the unit law $\mu_{DB} \circ D(\eta_{DB}) = \text{id}$, and the third holds by definition. Naturality follows from the naturality of q, μ and η . $\square$

9.3 The base-relative comonad

Definition 9.6. For a set Y put $C^B(Y) = Y \times A_B$, $C^B(f) = f \times \text{id}$, $\varepsilon_Y(y, c) = y$ and $\delta_Y(y, c) = ((y, c), c)$. Write C_{coh}^B for the same construction with A_B^{coh} .

This is the product (reader) comonad with parameter A_B , so the comonad laws are immediate. For $Y = B$ the underlying set $C^B(B) = B \times DB \times D^2B \times D^3B \times \text{Sh}(B)$ coincides with $C_3(B)$. The difference appears under the monad:

$$C^B(DY) = DY \times \underbrace{DB \times D^2B \times D^3B}_{\text{still three levels}} \times \text{Sh}(B).$$

9.4 The strict mixed distributive law

Definition 9.7. For $M = \sum_i a_i \delta_{(y_i, c_i)} \in DC^B(Y)$,

$$\lambda_Y^B(M) = \left(\sum_i a_i \delta_{y_i}, \sum_i a_i c_i \right) = (D(\text{pr}_1)M, \beta_B(D(\text{pr}_2)M)) \in C^B(DY),$$

where $\sum_i a_i c_i$ is the convex combination in A_B , taken level by level:

$$\sum_i a_i c_i = (\sum_i a_i \rho_i, \sum_i a_i \Pi_i, \sum_i a_i \Omega_i, \sum_i a_i \Sigma_i).$$

Theorem 9.8. $\lambda^B: \mathbf{DC}^B \Rightarrow \mathbf{C}^B \mathbf{D}$ is natural and satisfies all four axioms (U), (M), (Co), (Cm) of [definition 2.1](#). It restricts to a mixed distributive law $\mathbf{DC}_{\text{coh}}^B \Rightarrow \mathbf{C}_{\text{coh}}^B \mathbf{D}$.

Proof. Naturality. For $f: Y \rightarrow Y'$, $\mathbf{D}f \circ \mathbf{D}(\text{pr}_1) = \mathbf{D}(\text{pr}_1) \circ \mathbf{D}(f \times \text{id})$ and the context component is unchanged by $f \times \text{id}$.

(U). $\lambda^B(\delta_{(y,c)}) = (\delta_y, \beta_B(\delta_c)) = (\delta_y, c) = \mathbf{C}^B(\eta_Y)(y, c)$, by the unit law of the algebra β_B .

(Co). $\varepsilon_{\mathbf{D}Y} \circ \lambda_Y^B = \mathbf{D}(\text{pr}_1) = \mathbf{D}(\varepsilon_Y)$.

(M). First components: $\mathbf{D}(\text{pr}_1) \circ \mu_{\mathbf{C}^B Y} = \mu_Y \circ \mathbf{D}^2(\text{pr}_1)$ by naturality of μ , which is the first component of $\mathbf{C}^B(\mu_Y) \lambda_{\mathbf{D}Y}^B \mathbf{D}(\lambda_Y^B)$. Context components: the left side gives $\beta_B \circ \mathbf{D}(\text{pr}_2) \circ \mu = \beta_B \circ \mu_{AB} \circ \mathbf{D}^2(\text{pr}_2)$; the right side gives $\beta_B \circ \mathbf{D}(\beta_B \circ \mathbf{D}\text{pr}_2) = \beta_B \circ \mathbf{D}(\beta_B) \circ \mathbf{D}^2(\text{pr}_2)$, since $\mathbf{C}^B(\mu_Y)$ does not change the context. These agree by the associativity law $\beta_B \circ \mu_{AB} = \beta_B \circ \mathbf{D}(\beta_B)$ of the algebra.

(Cm). Let $\bar{c} = \beta_B(\mathbf{D}(\text{pr}_2)M)$. The left side is $\delta_{\mathbf{D}Y}(\lambda^B M) = (\lambda^B M, \bar{c})$. On the right, $\mathbf{D}(\delta_Y)M = \sum_i a_i \delta_{((y_i, c_i), c_i)}$; applying $\lambda_{\mathbf{C}^B Y}^B$ gives $(M, \bar{c})$, and applying $\mathbf{C}^B(\lambda_Y^B)$ gives $(\lambda^B M, \bar{c})$.

Restriction. β_B preserves A_B^{coh} by [lemma 9.3](#). $\square$

Remark 9.9. [Theorem 9.8](#) is an instance of the standard fact that a product comonad $- \times A$ distributes over a monad T whenever A carries a T -algebra structure. The content specific to this paper is the choice of A_B , the coherence sub-algebra and the canonical section σ_B . By the correspondence between mixed distributive laws and liftings [\[31\]](#), $\mathbf{D}$ lifts to the category of $\mathbf{C}^B$ -coalgebras.

9.5 Pooling, hierarchy and conditioning

Proposition 9.10 (Pooling commutes with hierarchy formation). *If $M = \sum_i a_i \delta_{(y_i, \sigma_B(\Sigma_i))}$, then $\lambda^B(M) = (\sum_i a_i \delta_{y_i}, \sigma_B(\sum_i a_i \Sigma_i))$.*

Proof. σ_B is affine ([proposition 9.5](#)). $\square$

Example 9.11 (Distinct posteriors survive at level two). Let $B = \{0, 1\}$, $\rho_0 = \delta_0$, $\rho_1 = \frac{1}{3}\delta_0 + \frac{2}{3}\delta_1$, and let $M = \frac{1}{2}\delta_{(y_0, \sigma(\delta_{h_0}))} + \frac{1}{2}\delta_{(y_1, \sigma(\delta_{h_1}))}$ with $q(h_k) = \rho_k$. The context of $\lambda^B(M)$ has level one $\frac{1}{2}\rho_0 + \frac{1}{2}\rho_1 = \frac{2}{3}\delta_0 + \frac{1}{3}\delta_1$ and level two $\frac{1}{2}\delta_{\rho_0} + \frac{1}{2}\delta_{\rho_1}$: the two posteriors are kept as distinct points of $\mathbf{DB}$, and level one is their barycenter, as coherence requires. In contrast, the averaging law of the absolute tower records level one through $\mathbf{D}(\eta)$ ([lemma 8.1](#)).

For evidence e with likelihood $L(e \mid \cdot)$ on $\text{Hyp}(B)$, conditioning a shadow is $\text{Cond}_e(\Sigma) = \mathcal{N}(L(e \mid \cdot) \Sigma)$ with evidence $Z_e(\Sigma) = \sum_h L(e \mid h) \Sigma(h)$.

Proposition 9.12 (Conditioning commutes with pooling). *Let $M = \sum_i a_i \delta_{(y_i, \sigma_B(\Sigma_i))}$ with $Z_i = Z_e(\Sigma_i) > 0$, and let $Z = \sum_i a_i Z_i$. Define the evidence-reweighted sample*

$$\text{Cond}_e(M) = \sum_i \frac{a_i Z_i}{Z} \delta_{(y_i, \sigma_B(\text{Cond}_e \Sigma_i))}.$$

Then the context of $\lambda^B(\text{Cond}_e M)$ is $\sigma_B(\text{Cond}_e(\sum_i a_i \Sigma_i))$.

Proof. $L \cdot \sum_i a_i \Sigma_i = \sum_i a_i Z_i \text{Cond}_e(\Sigma_i)$, so $\text{Cond}_e(\sum_i a_i \Sigma_i) = \sum_i (a_i Z_i / Z) \text{Cond}_e(\Sigma_i)$. Apply [proposition 9.10](#). $\square$

9.6 Change of base

Proposition 9.13. *For $f: B \rightarrow B'$ the map A_f is a morphism of D -algebras that preserves coherence and commutes with σ . The maps $\text{id}_Y \times A_f$ form a comonad morphism $C^B \Rightarrow C^{B'}$, and*

$$(\text{id}_{DY} \times A_f) \circ \lambda_Y^B = \lambda_Y^{B'} \circ D(\text{id}_Y \times A_f),$$

so λ^B is natural in the base.

Proof. A_f is a product of maps $D(g)$, hence affine, hence commutes with β . Coherence is preserved by naturality of μ and q ; compatibility with σ is [proposition 9.5](#). The comonad-morphism equations hold because δ copies the context and ε discards it. The displayed equation follows from $A_f \circ \beta_B = \beta_{B'} \circ D(A_f)$. $\square$

9.7 Comparison of the two options

At $Y = B$ the two towers have the same points. They differ in the comultiplication: the absolute comonad transports the context into distributions over C_3B through $\rho \mapsto (j_{B,c})_*\rho$, $\Pi \mapsto D^2(j_{B,c})\Pi$, $\Omega \mapsto D^3(j_{B,c})\Omega$, whereas the base-relative comonad copies it. That transport is what re-indexes the hierarchy and raises its order under D ; dropping it is what allows a strict law.

	levels at DY	(U)	(Co)	(M)	(Cm)
absolute, shift law λ^σ	four	no	yes	no	no
absolute, averaging law $\bar{\lambda}$	four	yes	yes	yes	no
base-relative, λ^B	three	yes	yes	yes	yes

Table 3: The two options. Only the base-relative tower keeps exactly three levels and admits a strict mixed distributive law ([theorem 9.8](#)).

10 Coherent lifting and obstruction certificates

The base-relative tower of [section 9](#) supports a precise notion of when a claimed posterior state is compatible with a hypothesis class at every level, and of what witnesses a failure. We fix a finite base B throughout.

10.1 Admissibility data and the admissible tower

Definition 10.1 (Admissibility data). Admissibility data consist of a finite set $S_1 \subseteq DB$ of *admissible posteriors* and a set $S_2 \subseteq D(S_1) \subseteq D^2B$ of *admissible second-order states* such that $\delta_\rho \in S_2$ for every $\rho \in S_1$. The *admissible tower* is

$$L_1 = \text{conv}(S_1) \subseteq DB, \quad L_2 = D(S_1) \subseteq D^2B, \quad L_3 = D(S_2) \subseteq D^3B.$$

Here $D(S)$ denotes the distributions supported in S .

The basic example comes from a finite admissible hypothesis class $H \subseteq \text{Hyp}(B)$, for instance the hypotheses consistent with the evidence: $S_1 = q_B(H)$. Two choices of S_2 bracket the possibilities. In the *committed regime* $S_2 = \{\delta_\rho : \rho \in S_1\}$: the only admissible second-order states are single admissible posteriors held with certainty. In the *inherited regime* $S_2 = D(S_1)$: every distribution over admissible posteriors is an admissible state. Richer hypothesis structure, for example composite hypotheses $h_1 \otimes h_2$ with an uncertain component, supplies intermediate choices.

Lemma 10.2. *Each L_k is convex, $\mu_B(L_2) = L_1$ and $\mu_{DB}(L_3) = L_2$. Moreover $\text{conv}(S_2) = L_2$.*

Proof. Convexity is clear. $\mu_B(\sum_\rho \Pi(\rho)\delta_\rho) = \sum_\rho \Pi(\rho)\rho$, so $\mu_B(D(S_1)) = \text{conv}(S_1)$. Since $S_2 \subseteq D(S_1)$ and $D(S_1)$ is convex, $\text{conv}(S_2) \subseteq D(S_1)$; conversely every $\Pi \in D(S_1)$ is $\sum_\rho \Pi(\rho)\delta_\rho$ with each $\delta_\rho \in S_2$. Finally $\mu_{DB}(D(S_2)) = \text{conv}(S_2) = L_2$. $\square$

Proposition 10.3 (Canonical contexts realize the committed tower). *Let $S_1 = q_B(H)$ and $S_2 = \{\delta_\rho : \rho \in S_1\}$. As Σ ranges over $D(H)$, the level components (ρ, Π, Ω) of the canonical contexts $\sigma_B(\Sigma)$ range exactly over the triples $(\mu_B\Pi, \Pi, D(\eta_{DB})\Pi)$ with $\Pi \in L_2$; in particular they fill L_1 , L_2 and L_3 .*

Proof. $\sigma_B(\Sigma)$ has $\Pi = D(q_B)\Sigma$, which ranges over $D(q_B(H)) = L_2$, and $\rho = \mu_B\Pi$, $\Omega = D(\eta_{DB})\Pi$. The last map sends $D(S_1)$ onto $D(\{\delta_\rho\}) = L_3$. $\square$

10.2 Claims, lifts and failure levels

Definition 10.4 (Claims and lifts). A *claim* is a partial tower $x = (x_1, x_2, x_3)$ with $x_k \in D^k B \cup \{*\}$ ($*$ meaning unspecified), consistent in the sense that $\mu_B(x_2) = x_1$ and $\mu_{DB}(x_3) = x_2$ whenever both sides are specified. A *coherent lift* of x is a triple $(\rho, \Pi, \Omega) \in L_1 \times L_2 \times L_3$ with $\mu_B\Pi = \rho$ and $\mu_{DB}\Omega = \Pi$ that agrees with x on every specified level. The *failure level* of x is

$$\varphi(x) = \min\{k : x_k \neq *, x_k \notin L_k\} \in \{1, 2, 3, \infty\}.$$

Theorem 10.5 (Lifting criterion). *Let x be a claim whose highest specified level is m .*

- (i) (Monotonicity) *If $x_k \notin L_k$ and x_j is specified for some $j > k$, then $x_j \notin L_j$. Hence a claim that fails at level k fails at every specified level above k .*
- (ii) (Criterion) *x has a coherent lift if and only if $x_m \in L_m$, if and only if $\varphi(x) = \infty$.*

Proof. (i) By lemma 10.2, $\mu(L_j) \subseteq L_{j-1}$, so if $x_j \in L_j$ then every specified x_i with $i < j$, being an iterated barycenter of x_j , lies in L_i . Contrapose.

(ii) If x has a lift, $x_m \in L_m$. Conversely suppose $x_m \in L_m$. Below m define the lift by barycenters: $\Pi = \mu_{DB}(x_3)$ if $m = 3$, and $\rho = \mu_B(\Pi)$; these lie in L_2, L_1 by lemma 10.2 and agree with any specified lower x_i by consistency. Above m use explicit sections: if $m \leq 2$ take $\Omega = D(\eta_{DB})(\Pi)$, which lies in $D(\{\delta_\rho : \rho \in S_1\}) \subseteq L_3$ and satisfies $\mu_{DB}\Omega = \Pi$; if $m = 1$ first write $x_1 = \sum_\rho w_\rho \rho$ with $\rho \in S_1$ and take $\Pi = \sum_\rho w_\rho \delta_\rho \in L_2$. The equivalence with $\varphi(x) = \infty$ follows from (i). $\square$

Example 10.6 (Each level can fail first). Let $B = \{0, 1\}$, $S_1 = \{\delta_0, \frac{1}{2}\delta_0 + \frac{1}{2}\delta_1\}$ and the committed regime $S_2 = \{\delta_{\delta_0}, \delta_{(1/2, 1/2)}\}$. Then $L_1 = \{(1-t, t) : 0 \leq t \leq \frac{1}{2}\}$.

- (a) *Level one.* $x_1 = \delta_1 \notin L_1$.
- (b) *Level two.* $x_2 = \delta_{(3/4, 1/4)}$. Its barycenter $(\frac{3}{4}, \frac{1}{4}) = \frac{1}{2}\delta_0 + \frac{1}{2}(\frac{1}{2}, \frac{1}{2})$ lies in L_1 , but $(\frac{3}{4}, \frac{1}{4}) \notin S_1$, so $x_2 \notin L_2$. The admissible $\frac{1}{2}\delta_{\delta_0} + \frac{1}{2}\delta_{(1/2, 1/2)}$ has the same barycenter.
- (c) *Level three.* $x_3 = \delta_{\Pi'}$ with $\Pi' = \frac{1}{2}\delta_{\delta_0} + \frac{1}{2}\delta_{(1/2, 1/2)}$. Then $\mu_{DB}x_3 = \Pi' \in L_2$, but $\Pi' \notin S_2$, so $x_3 \notin L_3$. The admissible $\frac{1}{2}\delta_{\delta_{\delta_0}} + \frac{1}{2}\delta_{\delta_{(1/2, 1/2)}}$ has the same barycenter.

In each case the failure is the same move one level up: a mixture of admissible states is asserted as a single state. At level two this is exactly the identity $\frac{1}{2}\delta_{\rho_1} + \frac{1}{2}\delta_{\rho_2} \neq \delta_{(\rho_1 + \rho_2)/2}$ behind the failure of the shift law (theorem 8.3).

Proposition 10.7 (Level-three obstructions need third-order data). *In the inherited regime $S_2 = D(S_1)$ one has $L_3 = \mu_{DB}^{-1}(L_2)$, so no claim fails first at level three.*

Proof. $\Omega \in D(D(S_1))$ iff every $\Pi' \in \text{supp } \Omega$ is supported in S_1 , iff

$$\text{supp } \mu_{DB}(\Omega) = \bigcup_{\Pi' \in \text{supp } \Omega} \text{supp } \Pi' \subseteq S_1. \quad \square$$

Thus hypothesis support alone produces obstructions only at levels one and two; a genuinely third-order obstruction requires a restriction on which second-order states are admissible.

10.3 Obstruction certificates

We measure obstructions by ℓ^1 distance, $d_k(x_k) = \min_{y \in L_k} \|x_k - y\|_1$, where for $k = 2, 3$ the norm is that of finitely supported measures on $D^{k-1}B$.

Theorem 10.8 (Level-one certificate). *$x_1 \notin L_1$ if and only if there is $f: B \rightarrow \mathbb{R}$ with*

$$\sum_b f(b)x_1(b) > \max_{\rho \in S_1} \sum_b f(b)\rho(b).$$

Quantitatively,

$$d_1(x_1) = \max_{\|f\|_\infty \leq 1} \left(\langle f, x_1 \rangle - \max_{\rho \in S_1} \langle f, \rho \rangle \right),$$

and any maximizing f is an optimal certificate.

Proof. L_1 is a polytope in $\mathbb{R}^B$, so the first statement is the separating hyperplane theorem (Farkas' lemma) [32]. For the identity, $d_1(x_1) = \min_{w \in \Delta_{S_1}} \max_{\|f\|_\infty \leq 1} \langle f, x_1 - \sum_\rho w_\rho \rho \rangle$, since $\max_{\|f\|_\infty \leq 1} \langle f, v \rangle = \|v\|_1$. The function is bilinear and both feasible sets are polytopes, so by the minimax theorem (equivalently LP duality [32]) the min and max may be exchanged, and $\min_w \langle f, \sum_\rho w_\rho \rho \rangle = \max_\rho \langle f, \rho \rangle$ is attained at a vertex. $\square$

Proposition 10.9 (Level-two and level-three certificates). *$x_2 \notin L_2$ iff some $\rho^* \in \text{supp } x_2$ lies outside S_1 , and $d_2(x_2) = 2x_2(DB \setminus S_1)$. Likewise $x_3 \notin L_3$ iff some $\Pi^* \in \text{supp } x_3$ lies outside S_2 , and $d_3(x_3) = 2x_3(D^2B \setminus S_2)$. The witnesses ρ^* and Π^* are certificates.*

Proof. For $y \in D(S)$ and outside mass $m = x(S^c)$, $\|x - y\|_1 \geq \sum_{S^c} x + |\sum_S (x - y)| = m + |(1 - m) - 1| = 2m$. Equality holds for y equal to x on S with the mass m moved to one point of S . $\square$

Proposition 10.10 (Pooling and obstructions). *For claims $x^{(i)}$ at the same level and weights a_i :*

(i) $d_1(\sum_i a_i x_1^{(i)}) \leq \sum_i a_i d_1(x_1^{(i)})$, and the inequality can be strict with every $x_1^{(i)} \notin L_1$ and the pooled claim in L_1 ;

(ii) $d_k(\sum_i a_i x_k^{(i)}) = \sum_i a_i d_k(x_k^{(i)})$ for $k = 2, 3$.

So pooling can hide level-one obstructions but never level-two or level-three ones.

Proof. (i) d_1 is the distance to a convex set, hence convex. For strictness take $S_1 = \{(\frac{1}{2}, \frac{1}{2})\}$: $d_1(\delta_0) = d_1(\delta_1) = 1$ while $\frac{1}{2}\delta_0 + \frac{1}{2}\delta_1 \in L_1$. (ii) Outside mass is linear in the measure; apply [proposition 10.9](#). In the same example, the level-two claim $\frac{1}{2}\delta_{\delta_0} + \frac{1}{2}\delta_{\delta_1}$ has $d_2 = 2$. $\square$

10.4 Effective admissibility and secondary obstructions

In [definition 10.1](#) every admissible posterior, held with certainty, is an admissible second-order state. This makes each projection $L_{k+1} \rightarrow L_k$ surjective ([lemma 10.2](#)), so a claim that passes the test at its own level always extends, and the tests of [theorem 10.5](#) are independent membership tests. A mechanism class need not have this property. A class that predicts a fair coin only by mixing two decisive hypotheses realizes the second-order state $\frac{1}{2}\delta_{\delta_0} + \frac{1}{2}\delta_{\delta_1}$ but not certainty in the uniform posterior, $\delta_{(1/2,1/2)}$. We now drop the requirement.

Definition 10.11 (General admissibility data and effective sets). *General admissibility data* are finite sets $S_1 \subseteq DB$ and $S_2 \subseteq D(S_1)$, with no further condition. They are *saturated* if $\delta_\rho \in S_2$ for every $\rho \in S_1$, as in [definition 10.1](#). The *local* sets L_1, L_2, L_3 are defined as in [definition 10.1](#), and the *effective* sets are

$$E_3 = L_3, \quad E_2 = \mu_{DB}(L_3) = \text{conv}(S_2), \quad E_1 = \mu_B(E_2) = \text{conv}(\mu_B(S_2)).$$

Claims and coherent lifts are defined as before, with lifts in $L_1 \times L_2 \times L_3$.

Lemma 10.12. *Each E_k is convex and $E_k \subseteq L_k$. The projections satisfy $\mu_B(L_2) = L_1$, $\mu_{DB}(L_3) = E_2 \subseteq L_2$ and $\mu_B(E_2) = E_1$. If the data are saturated, then $E_k = L_k$ for all k .*

Proof. $S_2 \subseteq D(S_1)$ and $D(S_1)$ is convex, so $E_2 \subseteq L_2$, and $E_1 = \mu_B(E_2) \subseteq \mu_B(L_2) = L_1$. The identities $\mu_B(D(S_1)) = \text{conv}(S_1)$ and $\mu_{DB}(D(S_2)) = \text{conv}(S_2)$ are proved as in [lemma 10.2](#), and $\mu_B(\text{conv } S_2) = \text{conv } \mu_B(S_2)$ because μ_B is affine. If the data are saturated, $\text{conv}(S_2) \supseteq \text{conv}\{\delta_\rho : \rho \in S_1\} = D(S_1)$. $\square$

Theorem 10.13 (Effective lifting criterion). *Let x be a claim whose highest specified level is m .*

- (i) *x has a coherent lift if and only if $x_m \in E_m$.*
- (ii) *If $x_m \in L_m \setminus E_m$, then every specified x_k lies in L_k , so x passes every local test of [theorem 10.5](#), yet x has no coherent lift.*
- (iii) (Certificates) *For $x_1 \in L_1$: $x_1 \notin E_1$ if and only if there is $f: B \rightarrow \mathbb{R}$ with $\langle f, x_1 \rangle > \max_{\Pi \in S_2} \langle f, \mu_B \Pi \rangle$, and the distance $d(x_1, E_1)$ satisfies the minimax identity of [theorem 10.8](#) with S_1 replaced by $\mu_B(S_2)$. For $x_2 \in L_2$: $x_2 \notin E_2$ if and only if there is $g: S_1 \rightarrow \mathbb{R}$ with*

$$\sum_\rho x_2(\rho)g(\rho) > \max_{\Pi \in S_2} \sum_\rho \Pi(\rho)g(\rho).$$

Proof. (i) If (ρ, Π, Ω) is a lift, then x_m is an iterated barycenter of $\Omega \in L_3$, so $x_m \in E_m$. Conversely, if $m = 3$ take $\Omega = x_3$. If $m \leq 2$, write $x_2 = \sum_j w_j \Pi_j$ (respectively $x_1 = \sum_j w_j \mu_B \Pi_j$) with $\Pi_j \in S_2$, and take $\Omega = \sum_j w_j \delta_{\Pi_j} \in L_3$, $\Pi = \mu_{DB} \Omega = \sum_j w_j \Pi_j \in L_2$ and $\rho = \mu_B \Pi \in L_1$. These agree with the specified levels of x by consistency. (ii) The lower levels of x are iterated barycenters of x_m , and $\mu(L_j) \subseteq L_{j-1}$ by [lemma 10.12](#). Combine with (i). (iii) E_1 and E_2 are convex hulls of finite sets in $\mathbb{R}^B$ and $\mathbb{R}^{S_1}$, so this is the separating hyperplane theorem; the minimax identity is proved as in [theorem 10.8](#). $\square$

The level-two certificate g is a test function on posteriors rather than on outcomes. It detects a claim whose first-order content may be entirely correct but whose second-order structure, meaning which posteriors it asserts and with what confidence, no admissible mechanism produces.

Definition 10.14 (Secondary obstruction). For $x_1 \in L_1$, the *fibre* of level-two lifts is $\mathcal{F}(x_1) = \{\Pi \in L_2 : \mu_B \Pi = x_1\}$, a nonempty compact convex polytope. For $\Pi \in L_2$, the obstruction to extending Π to level three is

$$o_3(\Pi) = \min_{\Pi' \in E_2} \|\Pi - \Pi'\|_1.$$

The *secondary obstruction* of x_1 is the set $I(x_1) = o_3(\mathcal{F}(x_1)) \subseteq [0, \infty)$.

Proposition 10.15 (Indeterminacy). *For $x_1 \in L_1$:*

- (i) $I(x_1)$ is a closed interval $[o_3^-(x_1), o_3^+(x_1)]$;
- (ii) x_1 has a coherent lift if and only if $0 \in I(x_1)$, if and only if $o_3^-(x_1) = 0$;
- (iii) $d(x_1, E_1) \leq o_3^-(x_1)$, where d is the ℓ^1 distance;
- (iv) o_3^- is convex on L_1 ;
- (v) if the data are saturated, then $I(x_1) = \{0\}$.

Proof. (i) $\mathcal{F}(x_1)$ is compact, convex and nonempty because $\mu_B(L_2) = L_1$, and o_3 is the distance to a convex set, hence continuous and convex. The image of a compact connected set under a continuous function is a compact interval. (ii) E_2 is closed, so $0 \in I(x_1)$ iff some $\Pi \in \mathcal{F}(x_1)$ lies in E_2 , iff (by the proof of [theorem 10.13](#)) x_1 has a coherent lift. (iii) μ_B is 1-Lipschitz for ℓ^1 : $\|\mu_B \Pi - \mu_B \Pi'\|_1 \leq \sum_{\rho} |\Pi(\rho) - \Pi'(\rho)| \|\rho\|_1 = \|\Pi - \Pi'\|_1$. Take $\Pi \in \mathcal{F}(x_1)$ attaining $o_3^-(x_1)$ and $\Pi' \in E_2$ nearest to it; then $\mu_B \Pi' \in E_1$. (iv) $(x_1, \Pi) \mapsto o_3(\Pi)$ restricted to the convex set $\{(x_1, \Pi) : \Pi \in L_2, \mu_B \Pi = x_1\}$ is jointly convex, and minimizing a jointly convex function over one variable gives a convex function of the other. (v) $E_2 = L_2$ by [lemma 10.12](#), so $o_3 \equiv 0$ on L_2 . $\square$

Example 10.16 (Gap, choice dependence and indeterminacy). Let $B = \{0, 1\}$, $u = \frac{1}{2}\delta_0 + \frac{1}{2}\delta_1$ and $\Pi_b = \frac{1}{2}\delta_{\delta_0} + \frac{1}{2}\delta_{\delta_1}$; in all three cases $S_2 = \{\Pi_b\}$, so $E_2 = \{\Pi_b\}$ and $E_1 = \{u\}$.

- (a) *Locally admissible, globally non-liftable.* With $S_1 = \{\delta_0, \delta_1\}$, $L_1 = DB$. The claim $x_1 = (\frac{1}{4}, \frac{3}{4})$ lies in L_1 , its only level-two lift is $\frac{1}{4}\delta_{\delta_0} + \frac{3}{4}\delta_{\delta_1}$, and $I(x_1) = \{\frac{1}{2}\} = \{d(x_1, E_1)\}$. The test function $f = (-1, 1)$ certifies $x_1 \notin E_1$ with margin $\frac{1}{2}$.
- (b) *Choice dependence.* With $S_1 = \{\delta_0, \delta_1, u\}$, the fibre over $x_1 = u$ is the segment $\Pi_t = \frac{1-t}{2}\delta_{\delta_0} + \frac{1+t}{2}\delta_{\delta_1} + t\delta_u$, $t \in [0, 1]$, and $o_3(\Pi_t) = 2t$. So $I(u) = [0, 2]$: the claim u lifts, through $\Pi_0 = \Pi_b$ and $\Omega = \delta_{\Pi_b}$, but the lift $\Pi_1 = \delta_u$ does not extend. As a level-two claim, δ_u passes every local test and is obstructed by the second-order certificate $g = \mathbf{1}_{\{u\}}$: $\langle g, \delta_u \rangle = 1 > 0 = \langle g, \Pi_b \rangle$.
- (c) *Nondegenerate indeterminacy without a lift.* With the same data and $x_1 = (\frac{3}{4}, \frac{1}{4})$, the fibre is $\{(\frac{3}{4} - \frac{c}{2})\delta_{\delta_0} + (\frac{1}{4} + \frac{c}{2})\delta_{\delta_1} + c\delta_u : 0 \leq c \leq \frac{1}{2}\}$, $o_3 = \frac{1}{2} + c$ along it, and $I(x_1) = [\frac{1}{2}, 1]$, with $o_3^-(x_1) = \frac{1}{2} = d(x_1, E_1)$.

Case (b) is the phenomenon described above: the first-order claim “fair coin” is correct, while the second-order claim “certainly the uniform posterior” is not, because the class produces fair-coin predictions only through uncertainty between decisive hypotheses. Pooling can hide the secondary obstruction, by [proposition 10.15](#)(iv): with $S_1 = \{\delta_0, \delta_1\}$, both δ_0 and δ_1 have $o_3^- = 1$, while their average u has $o_3^- = 0$.

Definition 10.17 (Stratified obstruction). Let x be a claim with highest specified level m . Its *obstruction type* is

- (a) *local at level k* if $\varphi(x) = k$, i.e. k is the least level with $x_k \notin L_k$;
- (b) *extension-obstructed* if $\varphi(x) = \infty$ but $x_m \notin E_m$; this happens only for $m \in \{1, 2\}$, and for $m = 1$ its secondary obstruction $I(x_1) \subseteq (0, \infty)$ records how far each choice of level-two lift is from extending;
- (c) *liftable* otherwise.

For a level-one claim this realizes the scheme of primary and secondary obstructions. Write $o_1(x_1) = d(x_1, L_1)$ for the primary obstruction. Once $o_1 = 0$, a level-two lift always exists, since $\mu_B(L_2) = L_1$, so the level-two obstruction vanishes automatically. The secondary obstruction is the set $I(x_1)$, which is defined only on the fibre of lifts and is said to vanish when $0 \in I(x_1)$. A level-one claim is therefore obstructed at level three precisely when $o_1(x_1) = 0$ and $0 \notin I(x_1)$. For saturated data every claim is local or liftable ([proposition 10.15\(v\)](#)), and [theorem 10.5](#) is recovered.

Remark 10.18 (Relation to classical obstruction theory). The structure parallels obstruction theory for lifting through a tower of fibrations: a primary obstruction, a secondary obstruction defined only once the primary one vanishes, dependence on the choice of intermediate lift, and an indeterminacy whose vanishing (here $0 \in I(x_1)$) is equivalent to the existence of a lift. The differences are substantial. Every level is a convex set, there are no homotopy groups, the obstructions are distances and separating functionals, and the indeterminacy is an interval of values, not a coset of a subgroup. We therefore speak of an *obstruction tower with indeterminacy*, not a Postnikov tower.

10.5 Interpretation and scope

Definition 10.19 (Obstructed claim). A claim is *obstructed* if it is not liftable in the sense of [definition 10.17](#). Its certificate is the test function of [theorem 10.8](#) or the witness state of [proposition 10.9](#) if it is local at level k , and the test function f or g of [theorem 10.13\(iii\)](#) if it is extension-obstructed.

We propose obstructed claims as a precise candidate for what is informally called a *hallucination*: a statement that may look admissible at a lower level but admits no coherent completion through the hierarchy generated by the admissible hypotheses. The definition is exact and computable, since every level reduces to a linear feasibility problem with a dual certificate, and [proposition 10.10](#) suggests testing at levels two and three, where aggregation cannot mask a local failure. The stratification of [definition 10.17](#) distinguishes a claim that is wrong at some level from one that is right at every level separately but cannot be completed coherently, and for level-one claims it records, through $I(x_1)$, how much the outcome depends on which intermediate explanation is chosen. Whether obstruction levels predict errors in trained models is an empirical question that this paper does not address.

Remark 10.20 (Why the obstructions are convex, not homotopical). Each $D^k B$ is convex and hence contractible, so homotopy-theoretic obstruction classes (such as k -invariants of a Postnikov tower) vanish identically on this tower, and μ has sections, so an unconstrained lift always exists. Obstructions arise only from the admissibility constraints, and they are convex-geometric: separating hyperplanes and support witnesses. The formal shape of primary and secondary obstructions survives nonetheless ([remark 10.18](#)).

Remark 10.21 (Relation to contextuality). When B is a product of outcome sets and claims are given as marginals on several contexts, the question whether they come from one admissible joint state is again a linear feasibility problem. This is the setting of sheaf-theoretic contextuality [1], where failure to admit a global section is the obstruction, and of its cohomological witnesses [2],

which are sufficient but not necessary for non-extendability. The single-context lifting problem of this section is the base case of that framework.

10.6 Sieving as iterated conditioning: a classical illustration

Analytic sieving gives a classical instance of the operations of this paper. We claim no new number theory; standard references are [18, 37]. Let P_0 be uniform on $\{1, \dots, N\}$ and, for a prime p , let $E_p = \{n : p \nmid n\}$.

- (a) *Conditioning.* Conditioning on an event E is the information projection [15]

$$P_0(\cdot \mid E) = \operatorname{argmin}\{\operatorname{KL}(P \parallel P_0) : P(E) = 1\}.$$

By the chain rule (section 4), conditioning successively on E_p for all primes $p < z$ equals conditioning once on their intersection, the z -rough integers.

- (b) *Where the process ends.* If $z^{k+1} > N$, every survivor $n > 1$ has at most k prime factors, and if $z > \sqrt{N}$ the survivors are exactly 1 and the primes in $[z, N]$. Integers with at most k prime factors are therefore intermediate stages, and the primes are the endpoint. For $z = N^{1/u}$ with $u > 1$ fixed, the number of survivors is asymptotic to $u\omega(u)N/\log N$, where ω is Buchstab's function [37, III.6]; Dickman's function governs the complementary smooth numbers.
- (c) *Parity as observational equivalence.* A linear sieve uses only the data $|\mathcal{A}_d| = \sum_{n \leq N, d \mid n} w(n)$ for squarefree d up to some level. Selberg observed that the weights $w_{\pm}(n) = 1 \pm \lambda(n)$, with λ Liouville's function, are supported on integers with an even, respectively odd, number of prime factors, while $|\mathcal{A}_d^{\pm}| = \lfloor N/d \rfloor \pm \lambda(d) \sum_{m \leq N/d} \lambda(m)$, and the last sum is $o(N/d)$. The two weightings are therefore indistinguishable by the sieve data up to errors that sieve methods must tolerate, although they separate parity completely [18, Ch. 16]. This is the situation of theorem 5.2 and proposition 5.7: two different states agree on every statistic the method observes.
- (d) *How the barrier is passed.* Chen's theorem that every large even integer is a prime plus a product of at most two primes [14] is consistent with the parity barrier. Passing the barrier, as for primes of the form $x^2 + y^4$ [17], required bilinear ("type II") estimates, that is, statistics beyond the sieve data. This matches proposition 5.7: identifiability is restored by observing more, not by a different geometry on the same observations.

11 Nondegeneracy and the likelihood-factorization criterion

Axiom 11.1 (Nondegeneracy). (a) $L(e \mid h_{\text{id}}) = 1$ for all e ; (b) $L(e \mid h) > 0$ for all h, e .

Remark 11.2 (Composite identifiability is too strong). The earlier version assumed that $\otimes$ is injective on pairs: $h_1 \otimes h_2 = h'_1 \otimes h'_2 \Rightarrow (h_1, h_2) = (h'_1, h'_2)$. Together with (A2) this forces $\operatorname{Hyp}(X) = \{h_{\text{id}}\}$: for any h , $h \otimes h_{\text{id}} = h = h_{\text{id}} \otimes h$, so $(h, h_{\text{id}}) = (h_{\text{id}}, h)$ and $h = h_{\text{id}}$. Restricting to supports does not help when the conclusion is quantified over all posteriors.

Theorem 11.3 (The "iff"). Assume axiom 11.1. Then $\operatorname{Cond}_e(P_1 \otimes P_2) = \operatorname{Cond}_e(P_1) \otimes \operatorname{Cond}_e(P_2)$ for all P_1, P_2 (whenever either side is defined) if and only if $L(e \mid h_1 \otimes h_2) = L(e \mid h_1)L(e \mid h_2)$ for all compatible pairs $h_1 \perp h_2$.

Proof. ($\Leftarrow$) is [proposition 4.12](#). ($\Rightarrow$) Let $h_1 \perp h_2$. If $h_1 = h_2 = h_{\text{id}}$ there is nothing to prove, since $L(e \mid h_{\text{id}}) = 1$. Otherwise fix $x \in X$ and $a, b \in (0, 1)$ and set $P_1 = a\delta_{(x, h_1)} + (1 - a)\delta_{(x, h_{\text{id}})}$, $P_2 = b\delta_{(x, h_2)} + (1 - b)\delta_{(x, h_{\text{id}})}$. Write $\alpha(h_1) = a$, $\alpha(h_{\text{id}}) = 1 - a$, $\beta(h_2) = b$, $\beta(h_{\text{id}}) = 1 - b$. By (A2) every pair from $\{h_1, h_{\text{id}}\} \times \{h_2, h_{\text{id}}\}$ is compatible, and since all mass sits on one outcome, $P_1 \otimes P_2$ has mass 1. For each $k \in \text{Hyp}(X)$ put

$$w(k) = \sum_{u \otimes v = k} \alpha(u)\beta(v) L(e \mid k), \quad w'(k) = \sum_{u \otimes v = k} \alpha(u)\beta(v) L(e \mid u)L(e \mid v),$$

summing over $(u, v) \in \{h_1, h_{\text{id}}\} \times \{h_2, h_{\text{id}}\}$. Then $\text{Cond}_e(P_1 \otimes P_2) = \mathcal{N}(w)$ and, using that all mass is on one outcome, $\text{Cond}_e P_1 \otimes \text{Cond}_e P_2 = \mathcal{N}(w')$. For pairs involving h_{id} , $L(e \mid u \otimes h_{\text{id}}) = L(e \mid u) = L(e \mid u)L(e \mid h_{\text{id}})$ by (A2) and [axiom 11.1\(a\)](#), so

$$w - w' = ab (L(e \mid h_1 \otimes h_2) - L(e \mid h_1)L(e \mid h_2)) \delta_{h_1 \otimes h_2}.$$

The composites $h_1 \otimes h_2$, h_1 , h_2 , h_{id} are not all equal (else $h_1 = h_2 = h_{\text{id}}$), so there is $k_0 \neq h_1 \otimes h_2$ among them, and $w(k_0) = w'(k_0) > 0$ by [axiom 11.1\(b\)](#). Equality $\mathcal{N}(w) = \mathcal{N}(w')$ means $w = cw'$ for some $c > 0$; evaluating at k_0 gives $c = 1$, hence $w = w'$ and $L(e \mid h_1 \otimes h_2) = L(e \mid h_1)L(e \mid h_2)$. $\square$

Remark 11.4. No identifiability or separation assumption is needed: the hypothesis h_{id} already supplies a reference point that fixes the normalizing constant. The earlier proof's constant $c = Z/(Z_1 Z_2)$ depended on the chosen posteriors and could not be set to 1 by a separate choice.

12 Coalgebras, bialgebras and the 2-categorical setting

This section describes the coalgebras, the lifted monad and the bialgebras of the base-relative law λ^B of [theorem 9.8](#), and makes precise which 2-category is involved.

12.1 The ambient 2-category

The monad D and the comonad C^B are a monad and a comonad on the object **Set** of the 2-category **Cat** of categories, functors and natural transformations, in the sense of the formal theory of monads [35]. Their structure maps $\eta, \mu, \varepsilon, \delta$ and the law λ^B are 2-cells of **Cat**, namely natural transformations between endofunctors of **Set**, and these 2-cells are not identities. [Theorem 9.8](#) is therefore a statement in **Cat**, and by the correspondence of [31] it is equivalent to a lifting of D to the category of C^B -coalgebras.

The categories of coalgebras, algebras and bialgebras obtained in this way are ordinary categories. Regarded as 2-categories (for instance by viewing each set as a discrete category), they are *locally discrete*: a natural transformation between functors of discrete categories exists only between equal functors, so the only 2-cells are identities. [Section 12.5](#) draws the consequences.

12.2 Coalgebras of the base-relative comonad

Proposition 12.1. *A map $\gamma: Y \rightarrow Y \times A_B$ is a C^B -coalgebra structure if and only if $\gamma = \langle \text{id}_Y, g \rangle$ for a (unique) map $g: Y \rightarrow A_B$, and a map $f: Y \rightarrow Y'$ is a coalgebra morphism $(Y, \langle \text{id}, g \rangle) \rightarrow (Y', \langle \text{id}, g' \rangle)$ if and only if $g' \circ f = g$. Hence $\text{Coalg}(C^B) \cong \mathbf{Set}/A_B$, and likewise $\text{Coalg}(C_{\text{coh}}^B) \cong \mathbf{Set}/A_B^{\text{coh}}$.*

Proof. The counit law $\varepsilon_Y \circ \gamma = \text{id}_Y$ forces the first component of γ to be the identity, so $\gamma = \langle \text{id}, g \rangle$. Coassociativity then holds automatically: both $C^B(\gamma) \circ \gamma$ and $\delta_Y \circ \gamma$ send y to $((y, g(y)), g(y))$. The condition $(f \times \text{id}) \circ \gamma = \gamma' \circ f$ reads $(f(y), g(y)) = (f(y), g'(f(y)))$. $\square$

Proposition 12.2 (Cofree coalgebras). *The forgetful functor $U: \text{Coalg}(\mathbf{C}^B) \rightarrow \mathbf{Set}$ has a right adjoint F , with $F(X) = (X \times A_B, \delta_X)$, i.e. $X \times A_B$ over A_B by pr_2 . Coalgebra morphisms $h: (Y, \gamma) \rightarrow F(X)$ correspond to maps $\varphi: Y \rightarrow X$ by $h \mapsto \text{pr}_1 \circ h$ and $\varphi \mapsto (\varphi \times \text{id}) \circ \gamma = \langle \varphi, g \rangle$. The counit of $U \dashv F$ at a set X is $\varepsilon_X = \text{pr}_1: X \times A_B \rightarrow X$, and the unit at a coalgebra (Y, γ) is the structure map $\gamma = \langle \text{id}_Y, g \rangle: (Y, \gamma) \rightarrow FU(Y, \gamma)$.*

Proof. Under [proposition 12.1](#), a morphism $(Y, g) \rightarrow (X \times A_B, \text{pr}_2)$ over A_B is a map $h = \langle \varphi, g \rangle$, determined by $\varphi = \text{pr}_1 \circ h$. The unit is the morphism corresponding to id_Y , which is $\langle \text{id}_Y, g \rangle$. $\square$

Remark 12.3. For a bare set X there is no natural map $X \rightarrow X \times A_B$: a map $x \mapsto (x, e)$ requires a choice of $e \in A_B$, and it is a coalgebra morphism only out of the constant coalgebra (X, const_e) . The unit of the cofree adjunction is the coalgebra structure itself, as in [proposition 12.2](#).

12.3 The lifted monad

Proposition 12.4. *Under the isomorphism of [proposition 12.1](#), the lifting of $\mathbf{D}$ determined by λ^B is*

$$\tilde{\mathbf{D}}(Y, g) = (\mathbf{D}Y, \beta_B \circ \mathbf{D}(g)), \quad \tilde{\mathbf{D}}(f) = \mathbf{D}(f),$$

with unit and multiplication those of $\mathbf{D}$. Conversely, λ_X^B is the coalgebra morphism $\tilde{\mathbf{D}}(FX) \rightarrow F(\mathbf{D}X)$ corresponding under [proposition 12.2](#) to $\mathbf{D}(\varepsilon_X) = \mathbf{D}(\text{pr}_1)$.

Proof. The lifting induced by a mixed law sends a coalgebra (Y, γ) to $(\mathbf{D}Y, \lambda_Y^B \circ \mathbf{D}\gamma)$ [\[31\]](#). For $\gamma = \langle \text{id}, g \rangle$ and $P \in \mathbf{D}Y$, $\lambda_Y^B(\mathbf{D}\gamma(P)) = (P, \beta_B(\mathbf{D}g(P)))$, which is the coalgebra $\langle \text{id}, \beta_B \circ \mathbf{D}g \rangle$. That η_Y and μ_Y are maps over A_B is the content of (U) and (M); directly, $\beta_B(\mathbf{D}g(\delta_y)) = g(y)$ by the unit law of the algebra β_B , and $\beta_B \circ \mathbf{D}g \circ \mu_Y = \beta_B \circ \mu_{A_B} \circ \mathbf{D}^2g = \beta_B \circ \mathbf{D}(\beta_B \circ \mathbf{D}g)$ by naturality of μ and the associativity law of β_B . For the converse, $\tilde{\mathbf{D}}(FX)$ is $\mathbf{D}(X \times A_B)$ over A_B by $\beta_B \circ \mathbf{D}(\text{pr}_2)$, so by [proposition 12.2](#) the morphism corresponding to $\mathbf{D}(\text{pr}_1)$ is $\langle \mathbf{D}(\text{pr}_1), \beta_B \circ \mathbf{D}(\text{pr}_2) \rangle = \lambda_X^B$. $\square$

12.4 Bialgebras

Recall that a λ -bialgebra is a set Y with a $\mathbf{D}$ -algebra structure $a: \mathbf{D}Y \rightarrow Y$ and a $\mathbf{C}^B$ -coalgebra structure γ such that $\gamma \circ a = \mathbf{C}^B(a) \circ \lambda_Y \circ \mathbf{D}(\gamma)$. Algebras of $\mathbf{D}$ are convex spaces [\[19\]](#), and their morphisms are affine maps.

Proposition 12.5. *A λ^B -bialgebra is exactly a convex space (Y, a) together with an affine map $g: (Y, a) \rightarrow (A_B, \beta_B)$. Hence the category of λ^B -bialgebras is isomorphic to the slice of the category of convex spaces over (A_B, β_B) , and the coherent variant to the slice over $(A_B^{\text{coh}}, \beta_B)$.*

Proof. With $\gamma = \langle \text{id}, g \rangle$ and $P \in \mathbf{D}Y$, the right side of the bialgebra condition is

$$\mathbf{C}^B(a)(P, \beta_B(\mathbf{D}g(P))) = (a(P), \beta_B(\mathbf{D}g(P))),$$

and the left side is $(a(P), g(a(P)))$. So the condition is $g \circ a = \beta_B \circ \mathbf{D}(g)$, which says that g is a morphism of $\mathbf{D}$ -algebras. $\square$

Example 12.6. By [proposition 9.5](#), $(\text{Sh}(B), \mu_{\text{Hyp}B}, \sigma_B)$ is a bialgebra with values in A_B^{coh} : the space of shadows, mixed as distributions, together with the canonical hierarchy it determines. For any coalgebra (Y, g) , the free bialgebra is $(\mathbf{D}Y, \mu_Y, \beta_B \circ \mathbf{D}g)$. A map that reads a context off a state non-affinely, for instance from the first hypothesis in the support of a shadow, violates the bialgebra condition.

So a bialgebra is a convex space of states with an affine readout of a three-level context: mixing states mixes their contexts level by level.

12.5 What two-dimensional monad theory does and does not add

Because the categories of [propositions 12.1, 12.4 and 12.5](#) are locally discrete as 2-categories, every 2-cell in them, invertible or not, is an identity. Consequently, for a 2-monad on such a 2-category, pseudo-algebras and pseudo-morphisms are strict ones, lax morphisms are strict ones, and the inclusion of strict algebras into pseudo-algebras is an isomorphism. The coherence and strictification results of Blackwell–Kelly–Power [\[6, 29\]](#), their notion of flexibility, two-dimensional monadicity [\[7\]](#) and doctrinal adjunction [\[28\]](#) therefore reduce to statements about ordinary categories and give no information beyond [theorem 9.8](#) and [propositions 12.1, 12.2, 12.4 and 12.5](#). We make no flexibility, strictification or 2-monadicity claims.

Genuine two-dimensional content would require three things that are not available in this paper:

- (i) an ambient 2-category in which contexts have non-identity morphisms, with specified hom-categories;
- (ii) a 2-monad extending D to that 2-category, with a specified action on 2-cells, and a comonad lifted along it;
- (iii) verification, for that 2-monad, of the codescent and doctrinal-adjunction hypotheses on which the results cited above depend.

12.6 Convex-order Beck comparisons for the absolute tower

Since the absolute tower admits no strict law ([theorems 8.3 and 8.7](#)), it is natural to ask whether an order on the relevant hom-sets turns its Beck defects into 2-cells. We test the most natural candidate, the convex order, and show that it fails in two independent ways. Throughout we write *convex-order Beck comparison* for an inequality between the two sides of a Beck diagram; no lax or oplax distributive law is claimed.

Definition 12.7 (Convex order). Let K be a convex subset of a real vector space. For finitely supported $\Pi, \Pi' \in D(K)$ write $\Pi \leq_{\text{cx}} \Pi'$ if $\sum_k \Pi(k)\varphi(k) \leq \sum_k \Pi'(k)\varphi(k)$ for every convex $\varphi: K \rightarrow \mathbb{R}$. Then Π and Π' have the same barycenter, and by Strassen’s theorem [\[34\]](#) $\Pi \leq_{\text{cx}} \Pi'$ if and only if there is a *martingale coupling*: a coupling π of Π and Π' such that, for each atom k of Π , the conditional distribution $\pi(k, \cdot)/\Pi(k)$ has barycenter k .

We apply this with $K = DX$ (a simplex in $\mathbb{R}^X$) and $K = C_3(DX) = DX \times E_3(DX)$, which is convex as a product of convex sets, each $D(Z)$ sitting affinely in the vector space of finitely supported functions on Z . Composites of convex functions with affine maps are convex.

Proposition 12.8 (Unit and multiplication comparisons). *For the shift law λ^σ of [theorem 8.3](#), on level-one components:*

- (a) (unit) the level one of $\lambda^\sigma \eta_{C_3}(x, c)$ is δ_ρ , that of $C_3(\eta_X)(x, c)$ is $D(\eta_X)\rho$, and $\delta_\rho \leq_{\text{cx}} D(\eta_X)\rho$;
- (b) (multiplication) with M' as in the proof of [theorem 8.3](#), the level one of $C_3(\mu_X)\lambda_{DX}^\sigma D(\lambda_X^\sigma)(M')$ is $\sum_\ell \nu_\ell \delta_{\bar{\rho}_\ell}$ with $\bar{\rho}_\ell = \sum_k a_{\ell k} \rho_{\ell k}$, that of $\lambda_X^\sigma \mu_{C_3 X}(M')$ is $\sum_{\ell, k} \nu_\ell a_{\ell k} \delta_{\rho_{\ell k}}$, and the former is $\leq_{\text{cx}}$ the latter.

Both inequalities are strict (not equalities) as soon as ρ is not a point mass, respectively some group $\{\rho_{\ell k}\}_k$ has two distinct members.

Proof. (a) The coupling $\pi(\rho, \delta_x) = \rho(x)$ is a martingale coupling, since $\sum_x \rho(x)\delta_x = \rho$. (b) The coupling $\pi(\bar{\rho}_\ell, \rho_{\ell k}) = \nu_\ell a_{\ell k}$ is a martingale coupling, since the conditional barycenter of $\bar{\rho}_\ell$ is $\sum_k a_{\ell k} \rho_{\ell k}$. Strictness: for the strictly convex $\varphi(v) = \|v\|_2^2$, Jensen's inequality is strict exactly when the averaged points are not all equal. $\square$

Proposition 12.9 (Orientation obstruction). *Write each Beck axiom of [definition 2.1](#) as an equation between its λ -first side ($\lambda \circ \eta_C$, respectively $\lambda \circ \mu_C$) and its other side ($C\eta$, respectively $C\mu \circ \lambda_T \circ T\lambda$). In [proposition 12.8](#) the λ -first side is below in the unit comparison and above in the multiplication comparison, and the reverse inequalities fail. Hence no choice of orientation of $\leq_{\text{cx}}$ makes both comparisons 2-cells of the same kind: reading $\leq_{\text{cx}}$ as a 2-cell, the unit comparison has the lax shape and the multiplication comparison the oplax shape, and reversing the order exchanges them.*

Proof. The orientations are those of [proposition 12.8](#); the reverse inequalities fail by its strictness statement. $\square$

The two comparisons are nonetheless consistent with each other: taking $M' = \delta_M$, the unit comparison at DX and monotonicity of $\text{D}(\mu_X)$ (pushforward along an affine map preserves $\leq_{\text{cx}}$) give $\text{C}_3(\mu_X)\lambda_{\text{DX}}^\sigma(\delta_{\lambda^\sigma M}) \leq_{\text{cx}} \lambda^\sigma(M)$ at level one, in agreement with [proposition 12.8\(b\)](#).

Proposition 12.10 (Comparability obstruction). *For both the shift law λ^σ and the averaging law $\bar{\lambda}$, take $X = \{0, 1\}$ and $M = \frac{1}{2}\delta_{(0, c_0)} + \frac{1}{2}\delta_{(0, c_1)}$ with c_i differing only in $\rho_i = \delta_i$. At level one, the two sides of $(\text{C}m)$ are*

$$L = \frac{1}{2}\delta_{(\delta_0, e)} + \frac{1}{2}\delta_{(\delta_1, e)}, \quad R = \frac{1}{2}\delta_{(\delta_0, e_0)} + \frac{1}{2}\delta_{(\delta_1, e_1)} \quad \text{in } \text{D}(\text{C}_3\text{DX}),$$

where e_i are the per-sample contexts and $e = \frac{1}{2}e_0 + \frac{1}{2}e_1$ is the pooled context. They have the same barycenter, and they are incomparable: $L \not\leq_{\text{cx}} R$ and $R \not\leq_{\text{cx}} L$.

Proof. The forms of L and R are computed in the proofs of [theorems 8.3](#) and [8.7](#). Let $s(P) = P(1)$ on DX and let $\tau(e')$ be the mass that the level-one component of e' assigns to δ_1 ; both are affine, and $\tau(e_0) = 0$, $\tau(e_1) = 1$, $\tau(e) = \frac{1}{2}$. The functions

$$\varphi_1(P, e') = (s(P) - \tau(e'))^2, \quad \varphi_2(P, e') = (\tau(e') - \frac{1}{2})^2$$

are convex. Then $\sum L\varphi_1 = \frac{1}{4} > 0 = \sum R\varphi_1$, so $L \not\leq_{\text{cx}} R$, and $\sum R\varphi_2 = \frac{1}{4} > 0 = \sum L\varphi_2$, so $R \not\leq_{\text{cx}} L$. Equivalently, in Strassen's terms: a martingale coupling must keep the first coordinate of each atom, which pins (δ_0, e) to (δ_0, e_0) and (δ_0, e_0) to (δ_0, e) , contradicting the second coordinate. $\square$

Corollary 12.11. *The naive convex-order relaxation of the absolute tower fails: the convex order orients the unit and multiplication defects of λ^σ oppositely ([proposition 12.9](#)) and cannot compare the comultiplication defect of either candidate law ([proposition 12.10](#)). In particular, neither λ^σ nor $\bar{\lambda}$ becomes a lax or oplax mixed distributive law by reading $\leq_{\text{cx}}$ as 2-cells.*

The first version of the paper proposed a “barycentric 2-cell” for the multiplication defect. [Proposition 12.8\(b\)](#) is the correct form of that idea, an order relation and not a map in **Set**, and [corollary 12.11](#) shows why it does not extend to a law.

Remark 12.12 (Toward a lifting order). The comparability obstruction has a clear source. The convex order compares only barycentric spreads, whereas the two sides of $(\text{C}m)$ differ in how contexts are attached to samples: one pooled context versus one context per sample, with equal barycenters. An enrichment that can orient this defect must compare the hierarchical coherence data of [section 10](#), namely whether and how a coherent lift through the tower exists, and not only barycenters. One

natural attempt fails and shows what is needed. Call $x \preceq x'$ if every admissibility datum (S_1, S_2) that lifts x' also lifts x . At level one this is equality: if $x \preceq x'$, take $S_1 = \{x'\}$, so $L_1 = \{x'\}$ and $x = x'$. A useful *lifting order* must therefore be relative to a fixed admissibility datum or carry additional structure, such as the certificates of [theorem 10.8](#) and [proposition 10.9](#), the effective sets of [definition 10.11](#), or the secondary obstruction of [definition 10.14](#). Constructing one that orients all three Beck defects of the absolute tower in the same direction is left open.

13 Empirical setup

The empirical setting is the CausalIDView benchmark of Jung, Seo, Byun, Lee and Song [27], which evaluates causal models across observational views; CausalIDView itself is not gradient-free. The method evaluated here is ours: a gradient-free treatment with structured shadows, developed in [5, Section 51]. The outcome space X consists of binary causal identification outcomes; descriptors θ are sparse multi-index hypotheses with spline links, nominally $r = 3$ directions; the posterior is sampled by gradient-free Metropolis. Targets are Manski-type partial-identification bounds [30] and instrumental-variable bounds.

On the 40-world benchmark ([appendix A](#)), the posterior under a generic prior with 16 pooled chains attains Manski RMSE 0.0934 and IV RMSE 0.1226, against 0.105 and 0.138 for the previous structured backbone. These numbers are close to the approximate TabPFN reference values [25]; since TabPFN was not rerun under the same protocol and no uncertainty intervals are reported, we do not claim parity. The flattening of the chain-pooling curve is consistent with, but does not by itself test, saturation of the level-two hierarchy. The 90% bootstrap intervals of the structured backbone cover the true bounds only 70–71% of the time, which indicates substantial miscalibration of that backbone. Coverage of the intervals produced by the generic-prior posterior has not been measured; bias-aware intervals remain open.

The categorical results above do not depend on these experiments, and the experiments do not test the categorical results.

14 Code and data availability

The verification scripts and related material are available at the repository below. The scripts are:

- `check_tower.py`: the absolute tower;
- `check_base_relative.py`: the base-relative tower;
- `check_lifting.py`: lifting and obstruction certificates;
- `check_obstruction_tower.py`: effective admissibility and secondary obstructions;
- `check_context_decomposition.py`: identifiability, contamination and sieving;
- `check_coalgebras.py`: coalgebras and bialgebras;
- `check_convex_order.py`: convex-order Beck comparisons;
- `attribution.py` (with `test_attribution.py` and `context_attribution/run_docs.py`): the sequential attribution of [section 5.4](#) and its experiment.

https://github.com/vkawasth/Trainingless-Transformer-NoGradientDescent/tree/main/Heirarchical_posteriors

15 Related work

Three-layer realisation, hierarchy and holonomy. [5, 3, 4].

Categorical probability. The finite distribution monad and Markov categories [13, 21]; a synthetic de Finetti theorem [20].

Distributive laws. Beck’s distributive laws [10]; the formal theory of monads [35]; mixed laws between a monad and a comonad and their correspondence with liftings [31]; weak distributive laws [36, 11, 12], with applications to the Vietoris monad [23] and to combining probabilistic and nondeterministic choice [24].

Convex duality. Farkas’ lemma, LP duality and the minimax identity used for obstruction certificates [32].

Information cohomology and contextuality. Baudot–Bennequin [9], Vigneaux [38], Abramsky–Brandenburger [1] and Abramsky–Mansfield–Barbosa [2].

Coalgebras and two-dimensional monad theory. Convex spaces as algebras of D [19]; two-dimensional monad theory [6, 29], two-dimensional monadicity [7] and doctrinal adjunction [28]; the convex order and Strassen’s theorem [34].

Attribution of generated text. Attribution of a model’s output to parts of its context by ablation and a surrogate model [42]; context-aware decoding, which contrasts the predictives with and without the context [47], the same pair as the context and the reference in section 5.4; EM, information projection and alternating minimisation [45, 15, 43]; likelihood ratios on the boundary and with nuisance parameters unidentified under the null [41, 46, 44]; misspecified likelihoods [48]; the corpora [40, 39].

Empirical. The CausalIDView benchmark [27]; the gradient-free method applied to it [5, Section 51]; TabPFN [25]; partial identification [30]; spiked random-matrix thresholds [8].

16 Summary of results

1. Base-relative tower: coherent contexts form a sub-algebra; the canonical hierarchy is an affine, natural section; λ^B is a strict mixed distributive law satisfying all four Beck axioms, natural in the base; pooling commutes with hierarchy formation and with conditioning (proved).
2. Coherent lifting: admissibility is convex at each level, failure is monotone, each level can fail first, level-three failures need third-order data, every failure has a certificate, and pooling hides level-one but not level-two or level-three obstructions (proved).
3. Obstruction tower with indeterminacy: for unsaturated data, lifts exist exactly on the effective sets, with first- and second-order certificates; the secondary obstruction on the fibre of level-two lifts depends on the chosen lift, ranges over an interval containing zero exactly when a lift exists, is bounded below by the effective distance, and can be hidden by pooling (proved, with worked examples).
4. Identifiability: source weights are unique for every mixture iff the sources are affinely independent; exclusive regions identify individual weights; supports-only information gives the core of a belief function, and for two sources the sharp interval $[x(R_A), 1 - x(R_B)]$; under unmodeled contamination the data certify a lower bound on the unmodeled weight, the distance test is

one-sided, and the number of unmodeled sources is not identifiable (proved). For token sequences the likelihood projection is unique under the sample form of affine independence, an absorbed reference loses the unexplained weight below an explicit fraction (proved), and blending against switching is a level-two distinction; the procedure is tested on real news text (table 1).

5. Coalgebras of the base-relative comonad are sets over A_B , the lifted monad is $(Y, g) \mapsto (DY, \beta_B \circ Dg)$, and bialgebras are convex spaces with an affine context map (proved). The ambient 2-category is **Cat**; the resulting categories are locally discrete, so two-dimensional monad theory adds nothing here. The naive convex-order relaxation of the absolute tower fails: unit and multiplication comparisons have opposite orientations and the comultiplication defect is incomparable (proved, with exact certificates).
6. The absolute tower grows to D^4X under D (remark 6.5).
7. $C_E(X) = X \times E(X)$ is a comonad for every E ; C_3 is an instance (proved).
8. (Hyp, q) with q natural; $J = \text{Hyp}(j)$ satisfies CL1–CL3; $K^{[k]}$ is natural and coheres with J via D^3j (proved).
9. Normalized partial algebra on $\text{Post}(X)$ with reweighted distributivity and conditioning laws (proved).
10. Affine first isomorphism theorem; congruence for $\otimes$ when $\perp$ is total; congruence for Conde iff $L(e \mid \cdot)$ is constant (proved, with counterexamples).
11. Absolute tower: shift law natural and (Co), fails (U), (M), (Cm); averaging law natural, (U), (Co), (M), fails (Cm). The unit axiom excludes point-mass hierarchy at level one (proved).
12. Composite identifiability trivializes Hyp; the “iff” holds under nondegeneracy alone (proved).

17 Open problems

1. For the absolute tower, does any mixed distributive law $DC_3 \Rightarrow C_3D$ satisfy all four axioms of definition 2.1? (For the base-relative tower this is settled by theorem 9.8.)
2. Bayesian updating as a morphism: make Conde of proposition 9.12 a natural operation on DC_{coh}^B compatible with λ^B , including evidence on outcomes.
3. Is there an enrichment in which all Beck defects of λ^σ or $\bar{\lambda}$ are witnessed by 2-cells of one orientation, giving a genuine lax or oplax law? The convex order does not work (corollary 12.11); remark 12.12 suggests comparing coherent lifts instead, and shows that the naive liftability preorder degenerates.
4. Specify a 2-category with non-identity 2-cells and a 2-monad extending D to it, as in section 12.5, and determine whether flexibility, strictification or 2-monadicity results then say anything about the coherent context algebra.
5. Weaken axiom 11.1 in theorem 11.3, e.g. allow zero likelihoods.
6. Secondary obstructions for higher-level claims and for restricted level-one data, and an efficient characterization of the upper end o_3^+ of the indeterminacy interval, which maximizes a convex function over the fibre.

7. Contamination at levels two and three: extend [theorem 5.10](#) to unmodeled second-order states, and characterize the least number of unmodeled sources when they are restricted to a family not closed under mixing.
8. Multi-context obstruction towers: combine [section 10.4](#) with the sheaf-theoretic setting of contextuality, where cohomological witnesses [2] could supply obstructions of a genuinely cohomological kind.
9. Approximate admissibility: replace exact support conditions at levels two and three by transport distances on DB and D^2B , whose Kantorovich duals would give graded certificates.
10. Empirical test: compute failure levels and certificates for model claims with known errors, and determine whether level-two and level-three obstructions predict them.
11. Sequential attribution ([section 5.4](#)) with a large language model as generator: calibration of the bands and of the unknown certificate when A_t , B_t , U_t are read from forward passes, more than two contexts, and windowed or change-point versions for long generations.
12. Extension to non-discrete outcome spaces (Giry monad).
13. Calibrated uncertainty and a test that could falsify the level-two saturation interpretation of the chain-pooling curve.

18 Conclusion

Keeping exactly three posterior levels requires fixing the base on which they live. In the absolute tower the levels are re-indexed by the object, so applying the monad raises the order to D^4X , and neither natural candidate law satisfies all Beck axioms; the unit axiom alone rules out point-mass hierarchy at level one. In the base-relative tower the levels stay over a fixed outcome base B , coherent contexts form a sub-algebra of the monad, the canonical hierarchy ρ_h , δ_{ρ_h} , $\delta_{\delta_{\rho_h}}$ is an affine section, and pooling defines a strict mixed distributive law satisfying all four Beck axioms. Pooling commutes with hierarchy formation and, after evidence reweighting, with conditioning, while distinct posteriors are retained at level two. The same tower supports an exact lifting problem: claims are checked level by level against the states a hypothesis class can produce, every failure comes with a certificate, and obstructions above level one survive aggregation. When the admissible second-order states are restricted, the lifting problem acquires a secondary obstruction that depends on the choice of intermediate lift, and a claim can be correct at every level separately yet admit no coherent completion. Attribution of a mixed context to its sources is exact for affinely independent sources, sharply bounded in the other regimes, and, under unmodeled contamination, complete only relative to a declared list of sources and a declared budget: the data certify how much is unexplained, but not how many unmodeled mechanisms there are. This makes “hallucination as obstructed lifting” a precise, testable definition. The coalgebras and bialgebras of the base-relative law are concrete, namely sets and convex spaces over the context algebra, and genuine two-dimensional structure is not present in this setting. For the absolute tower, relaxing the Beck equations to the convex order does not help: the unit and multiplication comparisons point in opposite directions and the comultiplication defect cannot be compared. The base-relative theorem therefore remains the control result, and the open question is which enrichment that compares coherent lifts, rather than barycentric spreads, could orient all the defects at once.

A Empirical tables

A.1 The 40-world benchmark

All numbers are on the 40-world set. The structured-backbone row uses 30 bootstrap refits per world; the posterior rows are computed from pooled Metropolis chains (8 or 16 chains), not from bootstrap refits. The protocols for the remaining rows are listed for confirmation in [appendix B](#).

Method	Manski RMSE	IV RMSE
Logistic regression (rebuild)	0.132	0.169
Structured backbone (previous best)	0.105	0.138
Posterior under generic prior, 8 chains	0.0937	0.1230
Posterior under generic prior, 16 chains	0.0934	0.1226
TabPFN, approximate reference [†]	≈ 0.09	≈ 0.12
Constant predictor	0.131	—

[†]Approximate values, not produced under this benchmark’s protocol; see [appendix B](#), item 13.

A.2 Chain-pooling curve

Chains	1	2	4	8	16
Manski	0.1005	0.0973	0.0953	0.0937	0.0934
IV	0.1319	0.1277	0.1250	0.1230	0.1226

Differences between 8 and 16 chains (about 3×10^{-4}) are reported without standard errors and should not be over-interpreted.

A.3 Coverage

The 90% bootstrap intervals of the structured backbone cover the true bounds 70–71% of the time (82% before the more regularized backbone). Coverage of the generic-prior posterior’s intervals has not been measured.

A.4 Shallow-MLP worlds

On shallow-MLP worlds the generic-prior posterior scores 0.104 against 0.103 for the structured backbone, i.e. no improvement. A misspecified level-two shadow is one possible explanation; it has not been tested.

A.5 Spectral pre-pass

The pre-pass recovers on average 2.5 of 6.6 planted directions and 4.5 of 12 active coordinates per direction. In spiked random-matrix models, the Baik–Ben Arous–Péché transition [8] gives a signal-strength threshold below which the top sample eigenvector carries no information about a planted direction. Whether our spectral pre-pass operates in such a regime has not been established, so we state the low recovery rate without attributing it to that threshold.

B Changes from earlier versions

Additions after version 8 (draft).

- (a) New [section 5.4](#): the likelihood projection of a token sequence onto declared contexts ([definition 5.14](#)), sample identifiability ([proposition 5.15](#)), the absorbed-reference proposition ([proposition 5.16](#)), profile bands with design-effect and boundary corrections, the unknown certificate, the level-two test of blending against switching, a decision rule, and an experiment on real news text ([table 1](#)).
- (b) New references [[39](#), [40](#), [41](#), [42](#), [43](#), [44](#), [45](#), [46](#), [47](#), [48](#)]; abstract, contributions, summary, related work, code availability and open problems updated.

Version 8.

- (a) New [section 5.2](#): identifiability of source weights ([proposition 5.7](#)), the supports-only interval and belief-function core ([proposition 5.8](#)), and a summary of the three regimes.
- (b) New [section 5.3](#): the contamination model, the certified lower bound on the unmodeled weight with its dual certificate and budgeted bounds ([theorem 5.10](#)), the one-sided distance test with a counterexample to its converse ([proposition 5.11](#)), and non-identifiability of the number of unmodeled sources ([proposition 5.12](#)).
- (c) New [section 10.6](#): analytic sieving as iterated conditioning, with the parity barrier as observational equivalence.
- (d) New references [[14](#), [15](#), [16](#), [17](#), [18](#), [22](#), [26](#), [33](#), [37](#)]; new script `check_context_decomposition.py`; abstract, introduction, summary, open problems, conclusion and checks updated.

Version 7.

- (a) New [section 10.4](#): general (unsaturated) admissibility data, effective sets, the effective lifting criterion with first- and second-order certificates ([theorem 10.13](#)), the secondary obstruction and its indeterminacy interval ([definition 10.14](#) and [proposition 10.15](#)), worked examples ([example 10.16](#)), the stratified obstruction types ([definition 10.17](#)), and the relation to classical obstruction theory ([remark 10.18](#)).
- (b) The definition of an obstructed claim now covers extension-obstructed claims; abstract, introduction, summary, open problems, conclusion and checks updated; new script `check_obstruction_tower.py`.

Version 6.

- (a) [Section 12.6](#) rewritten as convex-order Beck comparisons: unit and multiplication comparisons with exact martingale couplings ([proposition 12.8](#)), the orientation obstruction ([proposition 12.9](#)), the comparability obstruction for (Cm) for both candidate laws ([proposition 12.10](#)), and [corollary 12.11](#). The earlier suggestion that λ^σ might be a lax mixed distributive law in the convex order is withdrawn.
- (b) New [remark 12.12](#) on lifting orders, including the degeneracy of the naive liftability preorder; new “Structure of the argument” paragraph; abstract, summary, open problems and conclusion updated; new script `check_convex_order.py`.

Version 5. Relative to version 4 and to an intermediate draft that added two sections on 2-categorical lifting, BKP 2-monadicity and doctrinal adjunction:

- (a) The intermediate sections are replaced by [section 12](#). They worked in a 2-category of $\mathcal{C}^B$ -coalgebras whose 2-cells were described as natural transformations between unspecified “family actions”; no such structure had been defined, and on sets regarded as discrete categories the only 2-cells are identities. [Sections 12.1](#) and [12.5](#) state the ambient 2-category and what it does and does not provide.
- (b) The claimed cofree coalgebra used $\Delta^B(c) = (c, e_B)$ for a chosen e_B , which is not the comultiplication $(y, c) \mapsto ((y, c), c)$ of $\mathcal{C}^B$; the claimed adjunction unit $x \mapsto (x, e_B)$ is not the unit of the cofree adjunction and is not natural. Both are corrected in [proposition 12.2](#). The flexibility proof, which relied on them and used an undefined map κ_B in the role of the barycenter β_B , is removed, as are the strictification corollary, the 2-monadicity proposition (whose codescent and 2-cell arguments depended on the missing structure) and the doctrinal lax-structure formula.
- (c) Retained and proved: the lifted monad ([proposition 12.4](#)), with the coalgebra category identified as **Set**/ A_B ([proposition 12.1](#)), and the recovery of λ^B from the cofree coalgebra. New: bialgebras as convex spaces over A_B ([proposition 12.5](#)) and a first convex-order comparison of the shift law’s defects (superseded in version 6 by [section 12.6](#)).
- (d) References: the intermediate draft cited “Bourke and Lack, *Monads and theories*, Adv. Math. 287 (2015)”; *Monads and theories* is by Bourke and Garner, and two-dimensional monadicity is [7], now cited instead. A duplicate entry for [35] and an unused reference on weak bimonads are removed; broken cross-references are fixed.
- (e) CausalIDView is now attributed to [27]; it is not gradient-free, and the gradient-free method applied to it is cited to [5, Section 51].
- (f) Empirical corrections: the 70–71% coverage refers to the structured backbone’s bootstrap intervals, and coverage of the generic-prior posterior has not been measured; on shallow-MLP worlds 0.104 is the generic-prior posterior and 0.103 the structured backbone; the 30 bootstrap refits apply to the backbone row only, while posterior rows come from pooled Metropolis chains.

Version 4.

- (a) New [section 10](#): admissibility data, admissible tower, lifting criterion, failure levels, obstruction certificates, pooling behavior, and the interpretation of obstructed claims.
- (b) New reference [32]; abstract, introduction, summary, open problems, conclusion and computational checks updated.

Version 3.

- (a) New [section 9](#): base-relative tower with coherent contexts, canonical hierarchy section, and a strict mixed distributive law satisfying all four Beck axioms ([theorem 9.8](#)), natural in the base, with pooling and conditioning results.
- (b) New [remark 6.5](#) making explicit that $C_3(DX)$ contains D^4X , which the first version’s $\lambda^{[3]}$ omitted; the absolute-tower laws are now labelled Option 1.
- (c) Abstract, introduction, summary, open problems and conclusion rewritten around the design choice; repository link added.

Version 2 (relative to version 1).

1. Base category changed from finite sets to **Set**, since $D(X)$ is infinite.
2. Hyp redefined as $\Theta \times D(-)$; transport of directions and splines by $v_j \circ f$ removed (ill-typed, [remark 3.4](#)); the compatibility square is now naturality of q , not a definition.
3. Composition of posteriors normalized; distributivity and $\text{Cond}/\otimes$ commutation corrected ([propositions 4.8](#) and [4.12](#)); chain rule stated with its conditional likelihood.
4. Observational equivalence: the congruence for $\otimes$ needs total compatibility, not q -monoidality; the congruence for Cond claimed earlier is false and is replaced by [proposition 5.5](#).
5. The comonad is now an instance of a general theorem ([theorem 6.1](#)); Sh_3 replaced by Sh ; naturality of δ is proved.
6. J is defined as $\text{Hyp}(j)$ so that CL1–CL3 are proved; the seventh tuple component of the earlier construction is no longer needed.
7. K/J coherence corrected to use D^3j (it was ill-typed with D^2j).
8. $\lambda^{[3]}$ had four components instead of five and an ill-typed shadow; it is replaced by λ^σ . The earlier unit and comultiplication theorems were false ([theorem 8.3](#)). The averaging law $\bar{\lambda}$ and [lemma 8.1](#) are new.
9. The weak distributive law and 2-cell α are withdrawn ([remark 8.9](#)). The contradiction between the claimed proof of hexagonal coherence and its listing as an open problem is resolved by removal.
10. Composite identifiability (which trivializes Hyp) removed; it is not needed. The “iff” proof corrected ([theorem 11.3](#)); condition (c) of nondegeneracy (finiteness) dropped as automatic.
11. Bibliography corrected. The earlier references “Böhm and Lack, *Weak distributive laws*, J. Algebra 282 (2004)” and “Lack, *Composing monads using semimonads*, J. Algebra 281 (2004)” could not be located and are removed; they are replaced by [\[36, 11, 12, 23, 24, 31, 10\]](#). The de Finetti reference now cites its journal version [\[20\]](#).
12. Empirical claims toned down: no parity claim with TabPFN; BBP attribution removed; miscalibration flagged.

Items for the author to confirm.

13. Source and protocol of the TabPFN reference values in the benchmark table.
14. The nominal number of directions ($r = 3$ in the setup) versus the 6.6 planted directions in the spectral pre-pass.
15. Bootstrap standard errors for the chain-pooling curve.
16. Bibliographic details of the three self-citations [\[3, 4, 5\]](#), which are reproduced as given in the previous version.
17. Protocols for the logistic-regression and constant-predictor rows of the benchmark table.
18. The metric (Manski or IV RMSE) behind the shallow-MLP figures 0.104 and 0.103.

C Computational checks

All identities and counterexamples in [sections 5 and 8](#), [theorem 6.1](#) and [proposition 7.2](#) were checked in exact rational arithmetic on random instances with $X = \{0, 1\}$, two structural descriptors, and maps to a three-element set (one injective, one not). The checks confirm the comonad laws and naturality of ε, δ ; K/J coherence; naturality, (U), (Co) and (M) for $\bar{\lambda}$; the failure of (Cm) for $\bar{\lambda}$ on the explicit example of [theorem 8.7](#); naturality and (Co) for λ^σ and the failures of (U), (M), (Cm); and the counterexample of [proposition 5.5](#). These are in `check_tower.py`.

For the base-relative tower, `check_base_relative.py` checks, on random coherent and incoherent contexts over $B = \{0, 1\}$: the comonad laws; naturality of λ^B in Y ; all four axioms (U), (Co), (M), (Cm); closure of coherent contexts under barycenter; affineness and coherence of σ_B ; preservation of coherence and commutation with σ and λ^B under base change along an injective and a non-injective map; [example 9.11](#); and [proposition 9.12](#). All checks pass. Both scripts are in the repository cited in the code availability section.

For [section 10](#), `check_lifting.py` checks the three examples of [example 10.6](#) with their certificates; surjectivity $\mu(L_3) = L_2$ via explicit lifts; monotonicity of failure; the distance formula of [proposition 10.9](#) and its affinity under pooling; the minimax identity of [theorem 10.8](#) by solving the primal and dual linear programs; the cancellation example of [proposition 10.10](#); and that canonical contexts lie in the committed tower. The linear programs are solved in floating point; all other checks are exact. All 17 checks pass.

For [section 12](#), `check_coalgebras.py` checks that coalgebra structures are of the form $\langle \text{id}, g \rangle$ with coassociativity automatic; the cofree adjunction bijection; that η and μ are maps over A_B for the lifted monad; the recovery of λ^B from the cofree coalgebra; the bialgebra condition for $(\text{Sh}B, \mu, \sigma_B)$ and its failure for a non-affine context map; and the level-one convex-order inequalities on a family of convex test functions. All 10 checks pass.

For [section 12.6](#), `check_convex_order.py` certifies, in exact arithmetic: the unit and multiplication comparisons of [proposition 12.8](#) by constructing the martingale couplings and checking both marginals and the mean condition atom by atom, on random instances; the failure of the reverse inequalities, with the witness $\|v\|_2^2$; the resulting orientation obstruction of [proposition 12.9](#); and, for both λ^σ and $\bar{\lambda}$, that the two sides of (Cm) in [proposition 12.10](#) differ, have equal barycenters, and are incomparable, with the witnesses φ_1 and φ_2 . All 14 checks pass.

For [section 10.4](#), `check_obstruction_tower.py` verifies [example 10.16](#) exactly: the fibres, o_3 along each fibre, the indeterminacy intervals, the explicit lift and both certificates, and the pooling example. It checks exactly, on random instances, that μ_B is 1-Lipschitz, that points of E_1 have the explicit lift of [theorem 10.13](#), and that saturated data have $E_2 = L_2$. Using linear programs on random unsaturated data, it checks that $x_1 \in E_1$ iff $o_3^-(x_1) = 0$ and that $d(x_1, E_1) \leq o_3^-(x_1)$; for points outside E_1 , the dual certificate is rounded to rationals and verified to separate exactly. All 21 checks pass.

For [sections 5.2, 5.3 and 10.6](#), `check_context_decomposition.py` checks: exact and unique recovery of weights for affinely independent sources; for affinely dependent sources, an exact dependency giving two decompositions and a nondegenerate LP interval; the exclusive-region formula and per-observation attribution on a worked example; for supports only, that the interval of [proposition 5.8](#) is attained at both ends and cannot be extended, and, for three sources, that feasible allocations coincide with the core of the belief function; for contamination, exact primal-dual certificates bracketing $\mu_{\min}$, the interval $[\mu_{\min}, 1]$, the inequality $d \leq 2\mu_{\min}$, the equivalence $\mu_{\min} = 0 \iff x \in \text{conv}(A)$, the counterexample of [proposition 5.11](#), and the merging of unmodeled sources; and, for $N = 3000$, iterated versus single conditioning, the bounds on prime factors of survivors, the primes as endpoint, and the pure parity of Selberg's weights (with the relative

difference of their sieve data reported as an illustration). All 20 checks pass.

References

- [1] S. Abramsky and A. Brandenburger. The sheaf-theoretic structure of non-locality and contextuality. *New Journal of Physics*, 13:113036, 2011.
- [2] S. Abramsky, S. Mansfield, and R. S. Barbosa. The cohomology of non-locality and contextuality. *Electronic Proceedings in Theoretical Computer Science*, 95:1–14, 2012.
- [3] V. K. Awasthi. Recovering latent hierarchy without gradient descent. Technical report, Johns Hopkins University, September 2026. DOI: 10.13140/RG.2.2.24981.08167.
- [4] V. K. Awasthi. Holonomy obstructions for composed local operations. Technical report, Johns Hopkins University, September 2026. DOI: 10.13140/RG.2.2.25336.20483.
- [5] V. K. Awasthi. Three layers of realisation: hierarchy recovery, holonomy, information and contextuality. Technical report, Johns Hopkins University, September 2026. DOI: 10.13140/RG.2.2.24274.52163.
- [6] R. Blackwell, G. M. Kelly, and A. J. Power. Two-dimensional monad theory. *Journal of Pure and Applied Algebra*, 59(1):1–41, 1989.
- [7] J. Bourke. Two-dimensional monadicity. *Advances in Mathematics*, 252:708–747, 2014.
- [8] J. Baik, G. Ben Arous, and S. Péché. Phase transition of the largest eigenvalue for nonnull complex sample covariance matrices. *Annals of Probability*, 33(5):1643–1697, 2005.
- [9] P. Baudot and D. Bennequin. The homological nature of entropy. *Entropy*, 17(5):3253–3318, 2015.
- [10] J. Beck. Distributive laws. In *Seminar on Triples and Categorical Homology Theory*, volume 80 of *Lecture Notes in Mathematics*, pages 119–140. Springer, 1969.
- [11] G. Böhm. The weak theory of monads. *Advances in Mathematics*, 225(1):1–32, 2010. doi:10.1016/j.aim.2010.02.015.
- [12] G. Böhm, S. Lack, and R. Street. On the 2-categories of weak distributive laws. *Communications in Algebra*, 39(12):4567–4583, 2011. doi:10.1080/00927872.2011.616436.
- [13] T. Fritz. A synthetic approach to Markov kernels, conditional independence and theorems on sufficient statistics. *Advances in Mathematics*, 370:107239, 2020.
- [14] J. R. Chen. On the representation of a larger even integer as the sum of a prime and the product of at most two primes. *Scientia Sinica*, 16:157–176, 1973.
- [15] I. Csiszár. I -divergence geometry of probability distributions and minimization problems. *Annals of Probability*, 3(1):146–158, 1975.
- [16] A. P. Dempster. Upper and lower probabilities induced by a multivalued mapping. *Annals of Mathematical Statistics*, 38(2):325–339, 1967.

- [17] J. Friedlander and H. Iwaniec. The polynomial $X^2 + Y^4$ captures its primes. *Annals of Mathematics*, 148(3):945–1040, 1998.
- [18] J. Friedlander and H. Iwaniec. *Opera de Cribro*, volume 57 of *Colloquium Publications*. American Mathematical Society, 2010.
- [19] T. Fritz. Convex spaces I: definition and examples. arXiv:0903.5522, 2009.
- [20] T. Fritz, T. Gonda, and P. Perrone. De Finetti’s theorem in categorical probability. *Journal of Stochastic Analysis*, 2(4):Article 6, 2021. doi:10.31390/josa.2.4.06. arXiv:2105.02639.
- [21] T. Fritz, T. Gonda, P. Perrone, and E. F. Rischel. Representable Markov categories and comparison of statistical experiments in categorical probability. *Theoretical Computer Science*, 961:113896, 2023.
- [22] D. Gale. A theorem on flows in networks. *Pacific Journal of Mathematics*, 7(2):1073–1082, 1957.
- [23] R. Garner. The Vietoris monad and weak distributive laws. *Applied Categorical Structures*, 28(2):339–354, 2020. doi:10.1007/s10485-019-09582-w.
- [24] A. Goy and D. Petrişan. Combining probabilistic and non-deterministic choice via weak distributive laws. In *Proceedings of the 35th Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*, pages 454–464, 2020.
- [25] N. Hollmann, S. Müller, K. Eggenberger, and F. Hutter. TabPFN: A transformer that solves small tabular classification problems in a second. In *International Conference on Learning Representations (ICLR)*, 2023.
- [26] P. J. Huber. Robust estimation of a location parameter. *Annals of Mathematical Statistics*, 35(1):73–101, 1964.
- [27] H. Jung, G. Seo, H. Byun, J. Lee, and K. Song. What you observe determines how you identify causal effects: evaluating causal models across observational views. arXiv:2609.36881, 2026.
- [28] G. M. Kelly. Doctrinal adjunction. In *Category Seminar (Sydney, 1972/1973)*, volume 420 of *Lecture Notes in Mathematics*, pages 257–280. Springer, 1974.
- [29] S. Lack. Codescent objects and coherence. *Journal of Pure and Applied Algebra*, 175(1–3):223–241, 2002.
- [30] C. F. Manski. Nonparametric bounds on treatment effects. *American Economic Review*, 80(2):319–323, 1990.
- [31] J. Power and H. Watanabe. Combining a monad and a comonad. *Theoretical Computer Science*, 280(1–2):137–162, 2002. doi:10.1016/S0304-3975(01)00024-X.
- [32] A. Schrijver. *Theory of Linear and Integer Programming*. Wiley, Chichester, 1986.
- [33] G. Shafer. *A Mathematical Theory of Evidence*. Princeton University Press, 1976.
- [34] V. Strassen. The existence of probability measures with given marginals. *Annals of Mathematical Statistics*, 36(2):423–439, 1965.

- [35] R. Street. The formal theory of monads. *Journal of Pure and Applied Algebra*, 2(2):149–168, 1972.
- [36] R. Street. Weak distributive laws. *Theory and Applications of Categories*, 22(12):313–320, 2009.
- [37] G. Tenenbaum. *Introduction to Analytic and Probabilistic Number Theory*, 3rd edition, volume 163 of *Graduate Studies in Mathematics*. American Mathematical Society, 2015.
- [38] J. P. Vigneaux. *Topology of Statistical Systems: A Cohomological Approach to Information Theory*. PhD thesis, Université Paris Diderot, 2019.
- [39] T. Spinde, M. Plank, J.-D. Krieger, T. Ruas, B. Gipp, and A. Aizawa. Neural media bias detection using distant supervision with BABE – Bias Annotations By Experts. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pages 1166–1177, 2021.
- [40] L. Fan, M. White, E. Sharma, R. Su, P. K. Choubey, R. Huang, and L. Wang. In plain sight: media bias through the lens of factual reporting. In *Proceedings of EMNLP-IJCNLP*, pages 6343–6349, 2019.
- [41] H. Chernoff. On the distribution of the likelihood ratio. *Annals of Mathematical Statistics*, 25(3):573–578, 1954.
- [42] B. Cohen-Wang, H. Shah, K. Georgiev, and A. Madry. ContextCite: attributing model generation to context. In *Advances in Neural Information Processing Systems*, 2024.
- [43] I. Csiszár and G. Tusnády. Information geometry and alternating minimization procedures. *Statistics & Decisions*, Supplement Issue 1:205–237, 1984.
- [44] R. B. Davies. Hypothesis testing when a nuisance parameter is present only under the alternative. *Biometrika*, 64(2):247–254, 1977.
- [45] A. P. Dempster, N. M. Laird, and D. B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society, Series B*, 39(1):1–38, 1977.
- [46] S. G. Self and K.-Y. Liang. Asymptotic properties of maximum likelihood estimators and likelihood ratio tests under nonstandard conditions. *Journal of the American Statistical Association*, 82(398):605–610, 1987.
- [47] W. Shi, X. Han, M. Lewis, Y. Tsvetkov, L. Zettlemoyer, and W.-t. Yih. Trusting your evidence: hallucinate less with context-aware decoding. In *Proceedings of NAACL (Short Papers)*, 2024.
- [48] H. White. Maximum likelihood estimation of misspecified models. *Econometrica*, 50(1):1–25, 1982.