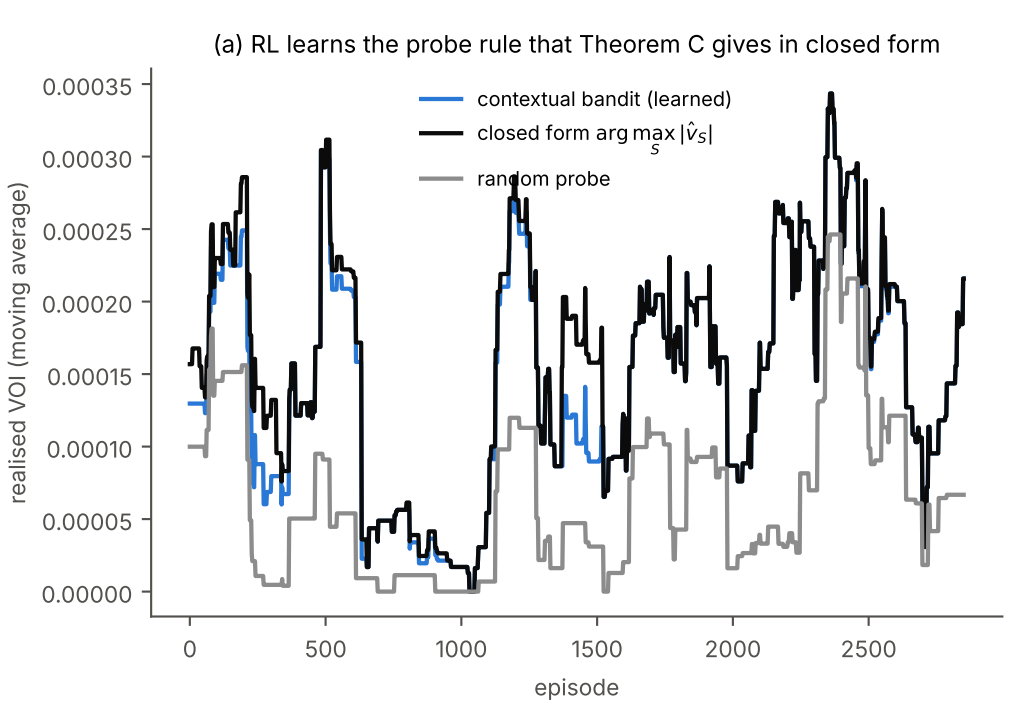
(a) RL learns the probe rule that Theorem C gives in closed form
contextual bandit (learned)
closed form arg max |v̂_s|
s
random probe
realised VOI (moving average)
0.00035
0.00030
0.00025
0.00020
0.00015
0.00010
0.00005
0.00000
0
500
1000
1500
2000
2500
episode

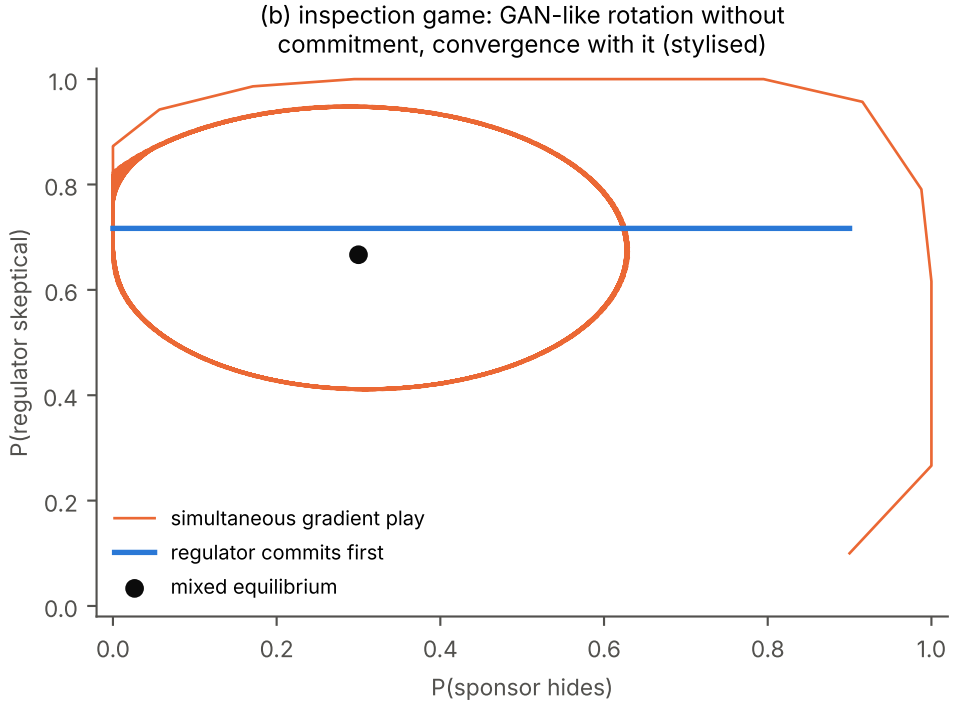
(b) inspection game: GAN-like rotation without commitment, convergence with it (stylised)
P(regulator skeptical)
1.0
0.8
0.6
0.4
0.2
0.0
simultaneous gradient play
regulator commits first
mixed equilibrium
0.0
0.2
0.4
0.6
0.8
1.0
P(sponsor hides)